

```
Using username "hadoop".
Authenticating with public key "imported-openssh-key"
Passphrase for key "imported-openssh-key":
```

```

  _  |  _  |  )
  _  |  (  _  /
  _  | \  _  |  _  |
  _  |  _  |  _  |

```

Amazon Linux AMI

```
https://aws.amazon.com/amazon-linux-ami/2018.03-release-notes/
9 package(s) needed for security, out of 17 available
Run "sudo yum update" to apply all updates.
```

EEEEEEEEEEEEEEEEEEEE	MMMMMMMM	MMMMMMMM	RRRRRRRRRRRRRRRR		
E::::::::::::::::::::E	M:::::::::M	M:::::::::M	R:::::::::::::::::R		
EE::::::::EEEEEEEEEE::E	M:::::::::M	M:::::::::M	R:::::RRRRRR:::::R		
E:::::E	EEEE	M:::::::::M	M:::::::::M	RR::::R	R:::::R
E:::::E	M:::::M:::M	M:::M:::::M	R:::R	R:::::R	
E:::::EEEEEEEEEE	M:::::M	M:::M	M:::M	M:::::M	R:::RRRRRR:::::R
E::::::::::::::::::::E	M:::::M	M:::M:::M	M:::::M	R:::::::::::::RR	
E:::::EEEEEEEEEE	M:::::M	M:::::M	M:::::M	R:::RRRRRR:::::R	
E:::::E	M:::::M	M:::M	M:::::M	R:::R	R:::::R
E:::::E	EEEE	MMM	M:::::M	R:::R	R:::::R
EE:::::EEEEEEEEEE::E	M:::::M	M:::::M	R:::R	R:::::R	
E::::::::::::::::::::E	M:::::M	M:::::M	RR::::R	R:::::R	
EEEEEEEEEEEEEEEEEEEE	MMMMMMMM	MMMMMMMM	RRRRRRR	RRRRRR	

```
[hadoop@ip-172-20-5-242 ~]$ ls
[hadoop@ip-172-20-5-242 ~]$ aws s3 cp s3://801198913bucket/FPGrowth.jar .
download: s3://801198913bucket/FPGrowth.jar to ./FPGrowth.jar
[hadoop@ip-172-20-5-242 ~]$ spark-submit --class org.FP.Driver
./FPGrowth.jar s3://801198913bucket/ExerciseData.txt 0.5 0.5
s3://801198913bucket/FPGrowthOutput
21/03/21 21:34:25 WARN SparkConf: The configuration key
'spark.kryoserializer.buffer.max.mb' has been deprecated as of Spark 1.4
and may be removed in the future. Please use the new key
'spark.kryoserializer.buffer.max' instead.
21/03/21 21:34:25 INFO SparkContext: Running Spark version 2.4.3
21/03/21 21:34:25 INFO SparkContext: Submitted application: Association
Action Rules Spark
21/03/21 21:34:25 INFO SecurityManager: Changing view acls to: hadoop
21/03/21 21:34:25 INFO SecurityManager: Changing modify acls to: hadoop
21/03/21 21:34:25 INFO SecurityManager: Changing view acls groups to:
21/03/21 21:34:25 INFO SecurityManager: Changing modify acls groups to:
21/03/21 21:34:25 INFO SecurityManager: SecurityManager: authentication
disabled; ui acls disabled; users with view permissions: Set(hadoop);
groups with view permissions: Set(); users with modify permissions:
Set(hadoop); groups with modify permissions: Set()
21/03/21 21:34:26 INFO Utils: Successfully started service 'sparkDriver'
on port 42931.
21/03/21 21:34:26 INFO SparkEnv: Registering MapOutputTracker
21/03/21 21:34:27 INFO SparkEnv: Registering BlockManagerMaster
21/03/21 21:34:27 INFO BlockManagerMasterEndpoint: Using
org.apache.spark.storage.DefaultTopologyMapper for getting topology
information
```

```
21/03/21 21:34:27 INFO BlockManagerMasterEndpoint:
BlockManagerMasterEndpoint up
21/03/21 21:34:27 INFO DiskBlockManager: Created local directory at
/mnt/tmp/blockmgr-dea45c0c-f0e2-4acc-bc7e-d76372ebe96b
21/03/21 21:34:27 INFO MemoryStore: MemoryStore started with capacity
424.4 MB
21/03/21 21:34:27 INFO SparkEnv: Registering OutputCommitCoordinator
21/03/21 21:34:28 INFO Utils: Successfully started service 'SparkUI' on
port 4040.
21/03/21 21:34:28 INFO SparkUI: Bound SparkUI to 0.0.0.0, and started at
http://ip-172-20-5-242.us-east-2.compute.internal:4040
21/03/21 21:34:28 INFO SparkContext: Added JAR
file:/home/hadoop/./FPGrowth.jar at spark://ip-172-20-5-242.us-east-
2.compute.internal:42931/jars/FPGrowth.jar with timestamp 1572212068400
21/03/21 21:34:28 INFO Executor: Starting executor ID driver on host
localhost
21/03/21 21:34:28 INFO Utils: Successfully started service
'org.apache.spark.network.netty.NettyBlockTransferService' on port 44177.
21/03/21 21:34:28 INFO NettyBlockTransferService: Server created on ip-
172-20-5-242.us-east-2.compute.internal:44177
21/03/21 21:34:28 INFO BlockManager: Using
org.apache.spark.storage.RandomBlockReplicationPolicy for block
replication policy
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager null re-registering with
master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager null
21/03/21 21:34:29 WARN Executor: Issue communicating with driver in
heartbeater
org.apache.spark.SparkException: Exception thrown in awaitResult:
    at
org.apache.spark.util.ThreadUtils$.awaitResult(ThreadUtils.scala:226)
    at
org.apache.spark.rpc.RpcTimeout.awaitResult(RpcTimeout.scala:75)
    at
org.apache.spark.rpc.RpcEndpointRef.askSync(RpcEndpointRef.scala:92)
    at
org.apache.spark.rpc.RpcEndpointRef.askSync(RpcEndpointRef.scala:76)
    at
org.apache.spark.storage.BlockManagerMaster.registerBlockManager(BlockMana
gerMaster.scala:73)
    at
org.apache.spark.storage.BlockManager.reregister(BlockManager.scala:339)
    at
org.apache.spark.executor.Executor.org$apache$spark$executor$Executor$$rep
ortHeartBeat(Executor.scala:845)
    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply$mcV$sp(Exe
cutor.scala:870)
    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply(Executor.s
cala:870)
```

```

    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply(Executor.s
cala:870)
    at
org.apache.spark.util.Utils$.logUncaughtExceptions(Utils.scala:1945)
    at
org.apache.spark.executor.Executor$$anon$2.run(Executor.scala:870)
    at
java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:511)
    at
java.util.concurrent.FutureTask.runAndReset(FutureTask.java:308)
    at
java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access
s$301(ScheduledThreadPoolExecutor.java:180)
    at
java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(S
cheduledThreadPoolExecutor.java:294)
    at
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:
1149)
    at
java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java
:624)
    at java.lang.Thread.run(Thread.java:748)
Caused by: java.lang.NullPointerException
    at
org.apache.spark.storage.BlockManagerMasterEndpoint.org$apache$spark$stora
ge$BlockManagerMasterEndpoint$$register(BlockManagerMasterEndpoint.scala:3
73)
    at
org.apache.spark.storage.BlockManagerMasterEndpoint$$anonfun$receiveAndRep
ly$1.applyOrElse(BlockManagerMasterEndpoint.scala:77)
    at
org.apache.spark.rpc.netty.Inbox$$anonfun$process$1.apply$mcV$sp(Inbox.sca
la:105)
    at org.apache.spark.rpc.netty.Inbox.safelyCall(Inbox.scala:205)
    at org.apache.spark.rpc.netty.Inbox.process(Inbox.scala:101)
    at
org.apache.spark.rpc.netty.Dispatcher$MessageLoop.run(Dispatcher.scala:221
)
    ... 3 more
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 ERROR Inbox: Ignoring error
java.lang.NullPointerException
    at
org.apache.spark.storage.BlockManagerMasterEndpoint.org$apache$spark$stora
ge$BlockManagerMasterEndpoint$$register(BlockManagerMasterEndpoint.scala:3
73)
    at
org.apache.spark.storage.BlockManagerMasterEndpoint$$anonfun$receiveAndRep
ly$1.applyOrElse(BlockManagerMasterEndpoint.scala:77)
    at
org.apache.spark.rpc.netty.Inbox$$anonfun$process$1.apply$mcV$sp(Inbox.sca
la:105)

```

```

        at org.apache.spark.rpc.netty.Inbox.safelyCall(Inbox.scala:205)
        at org.apache.spark.rpc.netty.Inbox.process(Inbox.scala:101)
        at
org.apache.spark.rpc.netty.Dispatcher$MessageLoop.run(Dispatcher.scala:221
)
        at
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:
1149)
        at
java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java
:624)
        at java.lang.Thread.run(Thread.java:748)
21/03/21 21:34:29 INFO BlockManager: BlockManager null re-registering with
master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager null
21/03/21 21:34:29 ERROR Inbox: Ignoring error
java.lang.NullPointerException
        at
org.apache.spark.storage.BlockManagerMasterEndpoint.org$apache$spark$stora
ge$BlockManagerMasterEndpoint$$register(BlockManagerMasterEndpoint.scala:3
73)
        at
org.apache.spark.storage.BlockManagerMasterEndpoint$$anonfun$receiveAndRep
ly$1.applyOrElse(BlockManagerMasterEndpoint.scala:77)
        at
org.apache.spark.rpc.netty.Inbox$$anonfun$process$1.apply$mcV$sp(Inbox.sca
la:105)
        at org.apache.spark.rpc.netty.Inbox.safelyCall(Inbox.scala:205)
        at org.apache.spark.rpc.netty.Inbox.process(Inbox.scala:101)
        at
org.apache.spark.rpc.netty.Dispatcher$MessageLoop.run(Dispatcher.scala:221
)
        at
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:
1149)
        at
java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java
:624)
        at java.lang.Thread.run(Thread.java:748)
21/03/21 21:34:29 WARN Executor: Issue communicating with driver in
heartbeater
org.apache.spark.SparkException: Exception thrown in awaitResult:
        at
org.apache.spark.util.ThreadUtils$.awaitResult(ThreadUtils.scala:226)
        at
org.apache.spark.rpc.RpcTimeout.awaitResult(RpcTimeout.scala:75)
        at
org.apache.spark.rpc.RpcEndpointRef.askSync(RpcEndpointRef.scala:92)
        at
org.apache.spark.rpc.RpcEndpointRef.askSync(RpcEndpointRef.scala:76)
        at
org.apache.spark.storage.BlockManagerMaster.registerBlockManager(BlockMana
gerMaster.scala:73)

```

```
    at
org.apache.spark.storage.BlockManager.reregister(BlockManager.scala:339)
    at
org.apache.spark.executor.Executor.org$apache$spark$executor$Executor$$reportHeartBeat(Executor.scala:845)
    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply$mcV$sp(Executor.scala:870)
    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply(Executor.scala:870)
    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply(Executor.scala:870)
    at
org.apache.spark.util.Utils$.logUncaughtExceptions(Utils.scala:1945)
    at
org.apache.spark.executor.Executor$$anon$2.run(Executor.scala:870)
    at
java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:511)
    at
java.util.concurrent.FutureTask.runAndReset(FutureTask.java:308)
    at
java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access$301(ScheduledThreadPoolExecutor.java:180)
    at
java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(ScheduledThreadPoolExecutor.java:294)
    at
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
    at
java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
    at java.lang.Thread.run(Thread.java:748)
Caused by: java.lang.NullPointerException
    at
org.apache.spark.storage.BlockManagerMasterEndpoint.org$apache$spark$storage$BlockManagerMasterEndpoint$$register(BlockManagerMasterEndpoint.scala:373)
    at
org.apache.spark.storage.BlockManagerMasterEndpoint$$anonfun$receiveAndReply$1.applyOrElse(BlockManagerMasterEndpoint.scala:77)
    at
org.apache.spark.rpc.netty.Inbox$$anonfun$process$1.apply$mcV$sp(Inbox.scala:105)
    at org.apache.spark.rpc.netty.Inbox.safelyCall(Inbox.scala:205)
    at org.apache.spark.rpc.netty.Inbox.process(Inbox.scala:101)
    at
org.apache.spark.rpc.netty.Dispatcher$MessageLoop.run(Dispatcher.scala:221)
)
... 3 more
```

21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManagerMasterEndpoint: Registering block
manager ip-172-20-5-242.us-east-2.compute.internal:44177 with 424.4 MB
RAM, BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal,
44177, None)
21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: external shuffle service port = 7337
21/03/21 21:34:29 INFO BlockManager: Initialized BlockManager:
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)

21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:29 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:29 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:29 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:29 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:30 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:30 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master

21/03/21 21:34:30 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:30 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:30 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:30 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:30 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:30 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:30 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:30 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:30 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:30 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:30 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:30 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:30 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:30 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:30 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:30 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:30 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:30 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:30 INFO Executor: Told to re-register on heartbeat

[illegible]

[illegible]

```

21/03/21 21:34:31 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:31 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:31 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:31 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:31 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:31 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:31 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:31 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:31 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:31 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:31 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:31 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:31 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:31 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:31 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:31 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)

```

21/03/21 21:34:31 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:31 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master

21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:32 INFO Executor: Told to re-register on heartbeat

21/03/21 21:34:32 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:32 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:32 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:32 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:33 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:33 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:33 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:33 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:33 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:33 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:33 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:33 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:33 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:33 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:33 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:33 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManager: Reporting 0 blocks to the master.

[illegible]

21/03/21 21:34:33 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:33 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:33 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:33 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:33 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)

21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master

21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:34 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:34 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:34 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:34 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:35 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:35 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:35 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:35 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:35 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:35 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:35 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering
with master
21/03/21 21:34:35 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:35 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177,
None)
21/03/21 21:34:35 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:35 INFO Executor: Told to re-register on heartbeat

[illegible]

21/03/21 21:34:35 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:35 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:35 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:35 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:35 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:35 INFO EventLoggingListener: Logging events to hdfs://var/log/spark/apps/local-1572212068609
21/03/21 21:34:35 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:35 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:35 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:35 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:35 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:36 INFO Executor: Told to re-register on heartbeat
21/03/21 21:34:36 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None) re-registering with master
21/03/21 21:34:36 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:36 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 44177, None)
21/03/21 21:34:36 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:34:38 INFO MemoryStore: Block broadcast_0 stored as values in memory (estimated size 237.8 KB, free 424.2 MB)
21/03/21 21:34:38 INFO MemoryStore: Block broadcast_0_piece0 stored as bytes in memory (estimated size 24.2 KB, free 424.2 MB)
21/03/21 21:34:38 INFO BlockManagerInfo: Added broadcast_0_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 24.2 KB, free: 424.4 MB)
21/03/21 21:34:38 INFO SparkContext: Created broadcast 0 from textFile at Driver.scala:18
21/03/21 21:34:39 WARN FPGrowth: Input data is not cached.
21/03/21 21:34:39 INFO GPLNativeCodeLoader: Loaded native gpl library
21/03/21 21:34:39 INFO LzoCodec: Successfully loaded & initialized native-lzo library [hadoop-lzo rev 7e6c862e89bc8db32c064454a55af74ddff73bae]
21/03/21 21:34:44 INFO FileInputFormat: Total input files to process : 1
21/03/21 21:34:44 INFO SparkContext: Starting job: count at FPGrowth.scala:217
21/03/21 21:34:44 INFO DAGScheduler: Got job 0 (count at FPGrowth.scala:217) with 1 output partitions

21/03/21 21:34:44 INFO DAGScheduler: Final stage: ResultStage 0 (count at FPGrowth.scala:217)
21/03/21 21:34:44 INFO DAGScheduler: Parents of final stage: List()
21/03/21 21:34:44 INFO DAGScheduler: Missing parents: List()
21/03/21 21:34:44 INFO DAGScheduler: Submitting ResultStage 0 (MapPartitionsRDD[2] at map at Driver.scala:20), which has no missing parents
21/03/21 21:34:45 INFO MemoryStore: Block broadcast_1 stored as values in memory (estimated size 3.4 KB, free 424.2 MB)
21/03/21 21:34:45 INFO MemoryStore: Block broadcast_1_piece0 stored as bytes in memory (estimated size 2.1 KB, free 424.2 MB)
21/03/21 21:34:45 INFO BlockManagerInfo: Added broadcast_1_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 2.1 KB, free: 424.4 MB)
21/03/21 21:34:45 INFO SparkContext: Created broadcast 1 from broadcast at DAGScheduler.scala:1201
21/03/21 21:34:45 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 0 (MapPartitionsRDD[2] at map at Driver.scala:20) (first 15 tasks are for partitions Vector(0))
21/03/21 21:34:45 INFO TaskSchedulerImpl: Adding task set 0.0 with 1 tasks
21/03/21 21:34:45 INFO TaskSetManager: Starting task 0.0 in stage 0.0 (TID 0, localhost, executor driver, partition 0, ANY, 7897 bytes)
21/03/21 21:34:45 INFO Executor: Running task 0.0 in stage 0.0 (TID 0)
21/03/21 21:34:45 INFO Executor: Fetching spark://ip-172-20-5-242.us-east-2.compute.internal:42931/jars/FPGrowth.jar with timestamp 1572212068400
21/03/21 21:34:45 INFO TransportClientFactory: Successfully created connection to ip-172-20-5-242.us-east-2.compute.internal/172.31.30.242:42931 after 104 ms (0 ms spent in bootstraps)
21/03/21 21:34:45 INFO Utils: Fetching spark://ip-172-20-5-242.us-east-2.compute.internal:42931/jars/FPGrowth.jar to /mnt/tmp/spark-be2b2756-6ddf-4cae-95ba-49e068570512/userFiles-e5c7a767-f0a1-412a-8250-9fda8e54af8b/fetchFileTemp991765250134519475.tmp
21/03/21 21:34:45 INFO Executor: Adding file:/mnt/tmp/spark-be2b2756-6ddf-4cae-95ba-49e068570512/userFiles-e5c7a767-f0a1-412a-8250-9fda8e54af8b/FPGrowth.jar to class loader
21/03/21 21:34:46 INFO HadoopRDD: Input split: s3://801198913bucket/ExerciseData.txt:0+72
21/03/21 21:34:46 INFO S3NativeFileSystem: Opening 's3://801198913bucket/ExerciseData.txt' for reading
21/03/21 21:34:46 INFO Executor: Finished task 0.0 in stage 0.0 (TID 0). 918 bytes result sent to driver
21/03/21 21:34:46 INFO TaskSetManager: Finished task 0.0 in stage 0.0 (TID 0) in 956 ms on localhost (executor driver) (1/1)
21/03/21 21:34:46 INFO TaskSchedulerImpl: Removed TaskSet 0.0, whose tasks have all completed, from pool
21/03/21 21:34:46 INFO DAGScheduler: ResultStage 0 (count at FPGrowth.scala:217) finished in 1.304 s
21/03/21 21:34:46 INFO DAGScheduler: Job 0 finished: count at FPGrowth.scala:217, took 1.612589 s
21/03/21 21:34:46 INFO BlockManagerInfo: Removed broadcast_1_piece0 on ip-172-20-5-242.us-east-2.compute.internal:44177 in memory (size: 2.1 KB, free: 424.4 MB)

21/03/21 21:34:46 INFO SparkContext: Starting job: collect at
FPGrowth.scala:257
21/03/21 21:34:47 INFO DAGScheduler: Registering RDD 4 (map at
FPGrowth.scala:254)
21/03/21 21:34:47 INFO DAGScheduler: Got job 1 (collect at
FPGrowth.scala:257) with 10 output partitions
21/03/21 21:34:47 INFO DAGScheduler: Final stage: ResultStage 2 (collect
at FPGrowth.scala:257)
21/03/21 21:34:47 INFO DAGScheduler: Parents of final stage:
List(ShuffleMapStage 1)
21/03/21 21:34:47 INFO DAGScheduler: Missing parents: List(ShuffleMapStage
1)
21/03/21 21:34:47 INFO DAGScheduler: Submitting ShuffleMapStage 1
(MapPartitionsRDD[4] at map at FPGrowth.scala:254), which has no missing
parents
21/03/21 21:34:47 INFO MemoryStore: Block broadcast_2 stored as values in
memory (estimated size 5.1 KB, free 424.2 MB)
21/03/21 21:34:47 INFO MemoryStore: Block broadcast_2_piece0 stored as
bytes in memory (estimated size 2.9 KB, free 424.2 MB)
21/03/21 21:34:47 INFO BlockManagerInfo: Added broadcast_2_piece0 in
memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 2.9 KB,
free: 424.4 MB)
21/03/21 21:34:47 INFO SparkContext: Created broadcast 2 from broadcast at
DAGScheduler.scala:1201
21/03/21 21:34:47 INFO DAGScheduler: Submitting 1 missing tasks from
ShuffleMapStage 1 (MapPartitionsRDD[4] at map at FPGrowth.scala:254)
(first 15 tasks are for partitions Vector(0))
21/03/21 21:34:47 INFO TaskSchedulerImpl: Adding task set 1.0 with 1 tasks
21/03/21 21:34:47 INFO TaskSetManager: Starting task 0.0 in stage 1.0 (TID
1, localhost, executor driver, partition 0, ANY, 7886 bytes)
21/03/21 21:34:47 INFO Executor: Running task 0.0 in stage 1.0 (TID 1)
21/03/21 21:34:47 INFO HadoopRDD: Input split:
s3://801198913bucket/ExerciseData.txt:0+72
21/03/21 21:34:47 INFO S3NativeFileSystem: Opening
's3://801198913bucket/ExerciseData.txt' for reading
21/03/21 21:34:47 INFO Executor: Finished task 0.0 in stage 1.0 (TID 1).
1193 bytes result sent to driver
21/03/21 21:34:47 INFO TaskSetManager: Finished task 0.0 in stage 1.0 (TID
1) in 221 ms on localhost (executor driver) (1/1)
21/03/21 21:34:47 INFO DAGScheduler: ShuffleMapStage 1 (map at
FPGrowth.scala:254) finished in 0.251 s
21/03/21 21:34:47 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:34:47 INFO DAGScheduler: running: Set()
21/03/21 21:34:47 INFO TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks
have all completed, from pool
21/03/21 21:34:47 INFO DAGScheduler: waiting: Set(ResultStage 2)
21/03/21 21:34:47 INFO DAGScheduler: failed: Set()
21/03/21 21:34:47 INFO DAGScheduler: Submitting ResultStage 2
(MapPartitionsRDD[6] at filter at FPGrowth.scala:256), which has no
missing parents
21/03/21 21:34:47 INFO MemoryStore: Block broadcast_3 stored as values in
memory (estimated size 3.8 KB, free 424.2 MB)
21/03/21 21:34:47 INFO MemoryStore: Block broadcast_3_piece0 stored as
bytes in memory (estimated size 2.3 KB, free 424.2 MB)

21/03/21 21:34:47 INFO BlockManagerInfo: Added broadcast_3_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 2.3 KB, free: 424.4 MB)

21/03/21 21:34:47 INFO SparkContext: Created broadcast 3 from broadcast at DAGScheduler.scala:1201

21/03/21 21:34:47 INFO DAGScheduler: Submitting 10 missing tasks from ResultStage 2 (MapPartitionsRDD[6] at filter at FPGrowth.scala:256) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9))

21/03/21 21:34:47 INFO TaskSchedulerImpl: Adding task set 2.0 with 10 tasks

21/03/21 21:34:47 INFO TaskSetManager: Starting task 2.0 in stage 2.0 (TID 2, localhost, executor driver, partition 2, PROCESS_LOCAL, 7662 bytes)

21/03/21 21:34:47 INFO Executor: Running task 2.0 in stage 2.0 (TID 2)

21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks

21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 16 ms

21/03/21 21:34:47 INFO Executor: Finished task 2.0 in stage 2.0 (TID 2). 1177 bytes result sent to driver

21/03/21 21:34:47 INFO TaskSetManager: Starting task 3.0 in stage 2.0 (TID 3, localhost, executor driver, partition 3, PROCESS_LOCAL, 7662 bytes)

21/03/21 21:34:47 INFO Executor: Running task 3.0 in stage 2.0 (TID 3)

21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks

21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 3 ms

21/03/21 21:34:47 INFO TaskSetManager: Finished task 2.0 in stage 2.0 (TID 2) in 83 ms on localhost (executor driver) (1/10)

21/03/21 21:34:47 INFO Executor: Finished task 3.0 in stage 2.0 (TID 3). 1134 bytes result sent to driver

21/03/21 21:34:47 INFO TaskSetManager: Starting task 4.0 in stage 2.0 (TID 4, localhost, executor driver, partition 4, PROCESS_LOCAL, 7662 bytes)

21/03/21 21:34:47 INFO Executor: Running task 4.0 in stage 2.0 (TID 4)

21/03/21 21:34:47 INFO TaskSetManager: Finished task 3.0 in stage 2.0 (TID 3) in 45 ms on localhost (executor driver) (2/10)

21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks

21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms

21/03/21 21:34:47 INFO Executor: Finished task 4.0 in stage 2.0 (TID 4). 1134 bytes result sent to driver

21/03/21 21:34:47 INFO TaskSetManager: Starting task 5.0 in stage 2.0 (TID 5, localhost, executor driver, partition 5, PROCESS_LOCAL, 7662 bytes)

21/03/21 21:34:47 INFO TaskSetManager: Finished task 4.0 in stage 2.0 (TID 4) in 33 ms on localhost (executor driver) (3/10)

21/03/21 21:34:47 INFO Executor: Running task 5.0 in stage 2.0 (TID 5)

21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks

21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 3 ms

21/03/21 21:34:47 INFO Executor: Finished task 5.0 in stage 2.0 (TID 5). 1134 bytes result sent to driver

21/03/21 21:34:47 INFO TaskSetManager: Starting task 6.0 in stage 2.0 (TID 6, localhost, executor driver, partition 6, PROCESS_LOCAL, 7662 bytes)

21/03/21 21:34:47 INFO TaskSetManager: Finished task 5.0 in stage 2.0 (TID 5) in 48 ms on localhost (executor driver) (4/10)
21/03/21 21:34:47 INFO Executor: Running task 6.0 in stage 2.0 (TID 6)
21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:47 INFO Executor: Finished task 6.0 in stage 2.0 (TID 6). 1134 bytes result sent to driver
21/03/21 21:34:47 INFO TaskSetManager: Starting task 0.0 in stage 2.0 (TID 7, localhost, executor driver, partition 0, ANY, 7662 bytes)
21/03/21 21:34:47 INFO TaskSetManager: Finished task 6.0 in stage 2.0 (TID 6) in 25 ms on localhost (executor driver) (5/10)
21/03/21 21:34:47 INFO Executor: Running task 0.0 in stage 2.0 (TID 7)
21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:47 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 4 ms
21/03/21 21:34:48 INFO Executor: Finished task 0.0 in stage 2.0 (TID 7). 1281 bytes result sent to driver
21/03/21 21:34:48 INFO TaskSetManager: Starting task 1.0 in stage 2.0 (TID 8, localhost, executor driver, partition 1, ANY, 7662 bytes)
21/03/21 21:34:48 INFO TaskSetManager: Finished task 0.0 in stage 2.0 (TID 7) in 65 ms on localhost (executor driver) (6/10)
21/03/21 21:34:48 INFO Executor: Running task 1.0 in stage 2.0 (TID 8)
21/03/21 21:34:48 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:48 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 7 ms
21/03/21 21:34:48 INFO Executor: Finished task 1.0 in stage 2.0 (TID 8). 1281 bytes result sent to driver
21/03/21 21:34:48 INFO TaskSetManager: Starting task 7.0 in stage 2.0 (TID 9, localhost, executor driver, partition 7, ANY, 7662 bytes)
21/03/21 21:34:48 INFO TaskSetManager: Finished task 1.0 in stage 2.0 (TID 8) in 46 ms on localhost (executor driver) (7/10)
21/03/21 21:34:48 INFO Executor: Running task 7.0 in stage 2.0 (TID 9)
21/03/21 21:34:48 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:48 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:34:48 INFO Executor: Finished task 7.0 in stage 2.0 (TID 9). 1281 bytes result sent to driver
21/03/21 21:34:48 INFO TaskSetManager: Starting task 8.0 in stage 2.0 (TID 10, localhost, executor driver, partition 8, ANY, 7662 bytes)
21/03/21 21:34:48 INFO TaskSetManager: Finished task 7.0 in stage 2.0 (TID 9) in 31 ms on localhost (executor driver) (8/10)
21/03/21 21:34:48 INFO Executor: Running task 8.0 in stage 2.0 (TID 10)
21/03/21 21:34:48 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:48 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:48 INFO Executor: Finished task 8.0 in stage 2.0 (TID 10). 1281 bytes result sent to driver

21/03/21 21:34:48 INFO TaskSetManager: Starting task 9.0 in stage 2.0 (TID 11, localhost, executor driver, partition 9, ANY, 7662 bytes)
21/03/21 21:34:48 INFO Executor: Running task 9.0 in stage 2.0 (TID 11)
21/03/21 21:34:48 INFO TaskSetManager: Finished task 8.0 in stage 2.0 (TID 10) in 21 ms on localhost (executor driver) (9/10)
21/03/21 21:34:48 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:48 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:34:48 INFO Executor: Finished task 9.0 in stage 2.0 (TID 11). 1324 bytes result sent to driver
21/03/21 21:34:48 INFO TaskSetManager: Finished task 9.0 in stage 2.0 (TID 11) in 13 ms on localhost (executor driver) (10/10)
21/03/21 21:34:48 INFO TaskSchedulerImpl: Removed TaskSet 2.0, whose tasks have all completed, from pool
21/03/21 21:34:48 INFO DAGScheduler: ResultStage 2 (collect at FPGrowth.scala:257) finished in 0.388 s
21/03/21 21:34:48 INFO DAGScheduler: Job 1 finished: collect at FPGrowth.scala:257, took 1.503713 s
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 71
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 62
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 38
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 46
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 44
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 33
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 74
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 25
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 49
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 52
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 31
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 29
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 27
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 61
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 39
21/03/21 21:34:48 INFO BlockManagerInfo: Removed broadcast_3_piece0 on ip-172-20-5-242.us-east-2.compute.internal:44177 in memory (size: 2.3 KB, free: 424.4 MB)
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 59
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 66
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 60
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 34
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 63
21/03/21 21:34:48 INFO ContextCleaner: Cleaned shuffle 0
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 54
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 50
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 37
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 28
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 36
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 64
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 55
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 56
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 42
21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 47

21/03/21 21:34:48 INFO BlockManagerInfo: Removed broadcast_2_piece0 on ip-172-20-5-242.us-east-2.compute.internal:44177 in memory (size: 2.9 KB, free: 424.4 MB)

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 73

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 72

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 57

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 51

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 69

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 30

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 32

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 70

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 40

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 58

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 45

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 35

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 53

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 67

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 26

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 68

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 48

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 65

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 41

21/03/21 21:34:48 INFO ContextCleaner: Cleaned accumulator 43

21/03/21 21:34:48 INFO deprecation: mapred.output.dir is deprecated. Instead, use mapreduce.output.fileoutputformat.outputdir

21/03/21 21:34:48 INFO HadoopMapRedCommitProtocol: Using output committer class org.apache.hadoop.mapred.DirectFileOutputCommitter

21/03/21 21:34:48 INFO DirectFileOutputCommitter: Nothing to setup since the outputs are written directly.

21/03/21 21:34:48 INFO SparkContext: Starting job: runJob at SparkHadoopWriter.scala:78

21/03/21 21:34:48 INFO DAGScheduler: Registering RDD 7 (flatMap at FPGrowth.scala:275)

21/03/21 21:34:48 INFO DAGScheduler: Registering RDD 12 (map at AssociationRules.scala:90)

21/03/21 21:34:48 INFO DAGScheduler: Registering RDD 11 (flatMap at AssociationRules.scala:78)

21/03/21 21:34:48 INFO DAGScheduler: Registering RDD 19 (repartition at Driver.scala:28)

21/03/21 21:34:48 INFO DAGScheduler: Got job 2 (runJob at SparkHadoopWriter.scala:78) with 1 output partitions

21/03/21 21:34:48 INFO DAGScheduler: Final stage: ResultStage 7 (runJob at SparkHadoopWriter.scala:78)

21/03/21 21:34:48 INFO DAGScheduler: Parents of final stage: List(ShuffleMapStage 6)

21/03/21 21:34:48 INFO DAGScheduler: Missing parents: List(ShuffleMapStage 6)

21/03/21 21:34:48 INFO DAGScheduler: Submitting ShuffleMapStage 3 (MapPartitionsRDD[7] at flatMap at FPGrowth.scala:275), which has no missing parents

21/03/21 21:34:48 INFO MemoryStore: Block broadcast_4 stored as values in memory (estimated size 7.1 KB, free 424.2 MB)

21/03/21 21:34:48 INFO MemoryStore: Block broadcast_4_piece0 stored as bytes in memory (estimated size 3.7 KB, free 424.2 MB)

21/03/21 21:34:48 INFO BlockManagerInfo: Added broadcast_4_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 3.7 KB, free: 424.4 MB)

21/03/21 21:34:48 INFO SparkContext: Created broadcast 4 from broadcast at DAGScheduler.scala:1201

21/03/21 21:34:48 INFO DAGScheduler: Submitting 1 missing tasks from ShuffleMapStage 3 (MapPartitionsRDD[7] at flatMap at FPGrowth.scala:275) (first 15 tasks are for partitions Vector(0))

21/03/21 21:34:48 INFO TaskSchedulerImpl: Adding task set 3.0 with 1 tasks

21/03/21 21:34:48 INFO TaskSetManager: Starting task 0.0 in stage 3.0 (TID 12, localhost, executor driver, partition 0, ANY, 7886 bytes)

21/03/21 21:34:48 INFO Executor: Running task 0.0 in stage 3.0 (TID 12)

21/03/21 21:34:48 INFO HadoopRDD: Input split: s3://801198913bucket/ExerciseData.txt:0+72

21/03/21 21:34:48 INFO S3NativeFileSystem: Opening 's3://801198913bucket/ExerciseData.txt' for reading

21/03/21 21:34:48 INFO Executor: Finished task 0.0 in stage 3.0 (TID 12). 1150 bytes result sent to driver

21/03/21 21:34:48 INFO TaskSetManager: Finished task 0.0 in stage 3.0 (TID 12) in 86 ms on localhost (executor driver) (1/1)

21/03/21 21:34:48 INFO TaskSchedulerImpl: Removed TaskSet 3.0, whose tasks have all completed, from pool

21/03/21 21:34:48 INFO DAGScheduler: ShuffleMapStage 3 (flatMap at FPGrowth.scala:275) finished in 0.110 s

21/03/21 21:34:48 INFO DAGScheduler: looking for newly runnable stages

21/03/21 21:34:48 INFO DAGScheduler: running: Set()

21/03/21 21:34:48 INFO DAGScheduler: waiting: Set(ShuffleMapStage 5, ShuffleMapStage 6, ResultStage 7, ShuffleMapStage 4)

21/03/21 21:34:48 INFO DAGScheduler: failed: Set()

21/03/21 21:34:48 INFO DAGScheduler: Submitting ShuffleMapStage 5 (MapPartitionsRDD[11] at flatMap at AssociationRules.scala:78), which has no missing parents

21/03/21 21:34:48 INFO MemoryStore: Block broadcast_5 stored as values in memory (estimated size 8.4 KB, free 424.2 MB)

21/03/21 21:34:49 INFO MemoryStore: Block broadcast_5_piece0 stored as bytes in memory (estimated size 4.2 KB, free 424.1 MB)

21/03/21 21:34:49 INFO BlockManagerInfo: Added broadcast_5_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 4.2 KB, free: 424.4 MB)

21/03/21 21:34:49 INFO SparkContext: Created broadcast 5 from broadcast at DAGScheduler.scala:1201

21/03/21 21:34:49 INFO DAGScheduler: Submitting 10 missing tasks from ShuffleMapStage 5 (MapPartitionsRDD[11] at flatMap at AssociationRules.scala:78) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9))

21/03/21 21:34:49 INFO TaskSchedulerImpl: Adding task set 5.0 with 10 tasks

21/03/21 21:34:49 INFO DAGScheduler: Submitting ShuffleMapStage 4 (MapPartitionsRDD[12] at map at AssociationRules.scala:90), which has no missing parents

21/03/21 21:34:49 INFO TaskSetManager: Starting task 5.0 in stage 5.0 (TID 13, localhost, executor driver, partition 5, PROCESS_LOCAL, 7651 bytes)

21/03/21 21:34:49 INFO Executor: Running task 5.0 in stage 5.0 (TID 13)

21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 5.0 in stage 5.0 (TID 13). 1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 6.0 in stage 5.0 (TID 14, localhost, executor driver, partition 6, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 5.0 in stage 5.0 (TID 13) in 18 ms on localhost (executor driver) (1/10)
21/03/21 21:34:49 INFO Executor: Running task 6.0 in stage 5.0 (TID 14)
21/03/21 21:34:49 INFO MemoryStore: Block broadcast_6 stored as values in memory (estimated size 8.4 KB, free 424.1 MB)
21/03/21 21:34:49 INFO MemoryStore: Block broadcast_6_piece0 stored as bytes in memory (estimated size 4.2 KB, free 424.1 MB)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 2 ms
21/03/21 21:34:49 INFO BlockManagerInfo: Added broadcast_6_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 4.2 KB, free: 424.4 MB)
21/03/21 21:34:49 INFO SparkContext: Created broadcast 6 from broadcast at DAGScheduler.scala:1201
21/03/21 21:34:49 INFO DAGScheduler: Submitting 10 missing tasks from ShuffleMapStage 4 (MapPartitionsRDD[12] at map at AssociationRules.scala:90) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9))
21/03/21 21:34:49 INFO TaskSchedulerImpl: Adding task set 4.0 with 10 tasks
21/03/21 21:34:49 INFO Executor: Finished task 6.0 in stage 5.0 (TID 14). 1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 5.0 in stage 4.0 (TID 15, localhost, executor driver, partition 5, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO Executor: Running task 5.0 in stage 4.0 (TID 15)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 6.0 in stage 5.0 (TID 14) in 32 ms on localhost (executor driver) (2/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:34:49 INFO Executor: Finished task 5.0 in stage 4.0 (TID 15). 1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 6.0 in stage 4.0 (TID 16, localhost, executor driver, partition 6, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 5.0 in stage 4.0 (TID 15) in 10 ms on localhost (executor driver) (1/10)
21/03/21 21:34:49 INFO Executor: Running task 6.0 in stage 4.0 (TID 16)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 7 ms
21/03/21 21:34:49 INFO Executor: Finished task 6.0 in stage 4.0 (TID 16). 1236 bytes result sent to driver

21/03/21 21:34:49 INFO TaskSetManager: Starting task 7.0 in stage 4.0 (TID 17, localhost, executor driver, partition 7, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 6.0 in stage 4.0 (TID 16) in 16 ms on localhost (executor driver) (2/10)
21/03/21 21:34:49 INFO Executor: Running task 7.0 in stage 4.0 (TID 17)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 7.0 in stage 4.0 (TID 17). 1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 8.0 in stage 4.0 (TID 18, localhost, executor driver, partition 8, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 7.0 in stage 4.0 (TID 17) in 21 ms on localhost (executor driver) (3/10)
21/03/21 21:34:49 INFO Executor: Running task 8.0 in stage 4.0 (TID 18)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 8.0 in stage 4.0 (TID 18). 1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 9.0 in stage 4.0 (TID 19, localhost, executor driver, partition 9, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO Executor: Running task 9.0 in stage 4.0 (TID 19)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 8.0 in stage 4.0 (TID 18) in 22 ms on localhost (executor driver) (4/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 4 ms
21/03/21 21:34:49 INFO Executor: Finished task 9.0 in stage 4.0 (TID 19). 1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 0.0 in stage 4.0 (TID 20, localhost, executor driver, partition 0, ANY, 7651 bytes)
21/03/21 21:34:49 INFO Executor: Running task 0.0 in stage 4.0 (TID 20)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 9.0 in stage 4.0 (TID 19) in 27 ms on localhost (executor driver) (5/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:34:49 INFO Executor: Finished task 0.0 in stage 4.0 (TID 20). 1365 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 1.0 in stage 4.0 (TID 21, localhost, executor driver, partition 1, ANY, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 0.0 in stage 4.0 (TID 20) in 82 ms on localhost (executor driver) (6/10)
21/03/21 21:34:49 INFO Executor: Running task 1.0 in stage 4.0 (TID 21)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 6 ms

21/03/21 21:34:49 INFO Executor: Finished task 1.0 in stage 4.0 (TID 21).
1365 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 2.0 in stage 4.0 (TID
22, localhost, executor driver, partition 2, ANY, 7651 bytes)
21/03/21 21:34:49 INFO Executor: Running task 2.0 in stage 4.0 (TID 22)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 1.0 in stage 4.0 (TID
21) in 30 ms on localhost (executor driver) (7/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 2.0 in stage 4.0 (TID 22).
1365 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 3.0 in stage 4.0 (TID
23, localhost, executor driver, partition 3, ANY, 7651 bytes)
21/03/21 21:34:49 INFO Executor: Running task 3.0 in stage 4.0 (TID 23)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 2.0 in stage 4.0 (TID
22) in 17 ms on localhost (executor driver) (8/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 4 ms
21/03/21 21:34:49 INFO Executor: Finished task 3.0 in stage 4.0 (TID 23).
1365 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 4.0 in stage 4.0 (TID
24, localhost, executor driver, partition 4, ANY, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 3.0 in stage 4.0 (TID
23) in 21 ms on localhost (executor driver) (9/10)
21/03/21 21:34:49 INFO Executor: Running task 4.0 in stage 4.0 (TID 24)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 2 ms
21/03/21 21:34:49 INFO Executor: Finished task 4.0 in stage 4.0 (TID 24).
1365 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 7.0 in stage 5.0 (TID
25, localhost, executor driver, partition 7, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 4.0 in stage 4.0 (TID
24) in 26 ms on localhost (executor driver) (10/10)
21/03/21 21:34:49 INFO TaskSchedulerImpl: Removed TaskSet 4.0, whose tasks
have all completed, from pool
21/03/21 21:34:49 INFO Executor: Running task 7.0 in stage 5.0 (TID 25)
21/03/21 21:34:49 INFO DAGScheduler: ShuffleMapStage 4 (map at
AssociationRules.scala:90) finished in 0.291 s
21/03/21 21:34:49 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:34:49 INFO DAGScheduler: running: Set(ShuffleMapStage 5)
21/03/21 21:34:49 INFO DAGScheduler: waiting: Set(ShuffleMapStage 6,
ResultStage 7)
21/03/21 21:34:49 INFO DAGScheduler: failed: Set()
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms

21/03/21 21:34:49 INFO Executor: Finished task 7.0 in stage 5.0 (TID 25).
1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 8.0 in stage 5.0 (TID
26, localhost, executor driver, partition 8, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO Executor: Running task 8.0 in stage 5.0 (TID 26)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 7.0 in stage 5.0 (TID
25) in 9 ms on localhost (executor driver) (3/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 3 ms
21/03/21 21:34:49 INFO Executor: Finished task 8.0 in stage 5.0 (TID 26).
1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 9.0 in stage 5.0 (TID
27, localhost, executor driver, partition 9, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 8.0 in stage 5.0 (TID
26) in 14 ms on localhost (executor driver) (4/10)
21/03/21 21:34:49 INFO Executor: Running task 9.0 in stage 5.0 (TID 27)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 9.0 in stage 5.0 (TID 27).
1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 0.0 in stage 5.0 (TID
28, localhost, executor driver, partition 0, ANY, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 9.0 in stage 5.0 (TID
27) in 12 ms on localhost (executor driver) (5/10)
21/03/21 21:34:49 INFO Executor: Running task 0.0 in stage 5.0 (TID 28)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 1 ms
21/03/21 21:34:49 INFO Executor: Finished task 0.0 in stage 5.0 (TID 28).
1279 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 1.0 in stage 5.0 (TID
29, localhost, executor driver, partition 1, ANY, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 0.0 in stage 5.0 (TID
28) in 23 ms on localhost (executor driver) (6/10)
21/03/21 21:34:49 INFO Executor: Running task 1.0 in stage 5.0 (TID 29)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 1.0 in stage 5.0 (TID 29).
1408 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 2.0 in stage 5.0 (TID
30, localhost, executor driver, partition 2, ANY, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 1.0 in stage 5.0 (TID
29) in 14 ms on localhost (executor driver) (7/10)
21/03/21 21:34:49 INFO Executor: Running task 2.0 in stage 5.0 (TID 30)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks

21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 2.0 in stage 5.0 (TID 30). 1365 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 3.0 in stage 5.0 (TID 31, localhost, executor driver, partition 3, ANY, 7651 bytes)
21/03/21 21:34:49 INFO Executor: Running task 3.0 in stage 5.0 (TID 31)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 2.0 in stage 5.0 (TID 30) in 20 ms on localhost (executor driver) (8/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 3.0 in stage 5.0 (TID 31). 1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 4.0 in stage 5.0 (TID 32, localhost, executor driver, partition 4, ANY, 7651 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 3.0 in stage 5.0 (TID 31) in 8 ms on localhost (executor driver) (9/10)
21/03/21 21:34:49 INFO Executor: Running task 4.0 in stage 5.0 (TID 32)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 4.0 in stage 5.0 (TID 32). 1236 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Finished task 4.0 in stage 5.0 (TID 32) in 13 ms on localhost (executor driver) (10/10)
21/03/21 21:34:49 INFO TaskSchedulerImpl: Removed TaskSet 5.0, whose tasks have all completed, from pool
21/03/21 21:34:49 INFO DAGScheduler: ShuffleMapStage 5 (flatMap at AssociationRules.scala:78) finished in 0.449 s
21/03/21 21:34:49 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:34:49 INFO DAGScheduler: running: Set()
21/03/21 21:34:49 INFO DAGScheduler: waiting: Set(ShuffleMapStage 6, ResultStage 7)
21/03/21 21:34:49 INFO DAGScheduler: failed: Set()
21/03/21 21:34:49 INFO DAGScheduler: Submitting ShuffleMapStage 6 (MapPartitionsRDD[19] at repartition at Driver.scala:28), which has no missing parents
21/03/21 21:34:49 INFO MemoryStore: Block broadcast_7 stored as values in memory (estimated size 5.5 KB, free 424.1 MB)
21/03/21 21:34:49 INFO MemoryStore: Block broadcast_7_piece0 stored as bytes in memory (estimated size 3.0 KB, free 424.1 MB)
21/03/21 21:34:49 INFO BlockManagerInfo: Added broadcast_7_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 3.0 KB, free: 424.4 MB)
21/03/21 21:34:49 INFO SparkContext: Created broadcast 7 from broadcast at DAGScheduler.scala:1201
21/03/21 21:34:49 INFO DAGScheduler: Submitting 10 missing tasks from ShuffleMapStage 6 (MapPartitionsRDD[19] at repartition at Driver.scala:28) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9))
21/03/21 21:34:49 INFO TaskSchedulerImpl: Adding task set 6.0 with 10 tasks

21/03/21 21:34:49 INFO TaskSetManager: Starting task 0.0 in stage 6.0 (TID 33, localhost, executor driver, partition 0, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO Executor: Running task 0.0 in stage 6.0 (TID 33)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 0.0 in stage 6.0 (TID 33). 1270 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 1.0 in stage 6.0 (TID 34, localhost, executor driver, partition 1, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO Executor: Running task 1.0 in stage 6.0 (TID 34)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 0.0 in stage 6.0 (TID 33) in 58 ms on localhost (executor driver) (1/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 1.0 in stage 6.0 (TID 34). 1227 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 2.0 in stage 6.0 (TID 35, localhost, executor driver, partition 2, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 1.0 in stage 6.0 (TID 34) in 11 ms on localhost (executor driver) (2/10)
21/03/21 21:34:49 INFO Executor: Running task 2.0 in stage 6.0 (TID 35)
21/03/21 21:34:49 INFO BlockManagerInfo: Removed broadcast_5_piece0 on ip-172-20-5-242.us-east-2.compute.internal:44177 in memory (size: 4.2 KB, free: 424.4 MB)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 2.0 in stage 6.0 (TID 35). 1227 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 3.0 in stage 6.0 (TID 36, localhost, executor driver, partition 3, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 2.0 in stage 6.0 (TID 35) in 12 ms on localhost (executor driver) (3/10)
21/03/21 21:34:49 INFO Executor: Running task 3.0 in stage 6.0 (TID 36)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks

21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 2 non-empty blocks including 2 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:34:49 INFO Executor: Finished task 3.0 in stage 6.0 (TID 36). 1227 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 4.0 in stage 6.0 (TID 37, localhost, executor driver, partition 4, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 3.0 in stage 6.0 (TID 36) in 14 ms on localhost (executor driver) (4/10)
21/03/21 21:34:49 INFO Executor: Running task 4.0 in stage 6.0 (TID 37)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:34:49 INFO Executor: Finished task 4.0 in stage 6.0 (TID 37). 1227 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 5.0 in stage 6.0 (TID 38, localhost, executor driver, partition 5, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 4.0 in stage 6.0 (TID 37) in 20 ms on localhost (executor driver) (5/10)
21/03/21 21:34:49 INFO Executor: Running task 5.0 in stage 6.0 (TID 38)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO BlockManagerInfo: Removed broadcast_4_piece0 on ip-172-20-5-242.us-east-2.compute.internal:44177 in memory (size: 3.7 KB, free: 424.4 MB)
21/03/21 21:34:49 INFO Executor: Finished task 5.0 in stage 6.0 (TID 38). 1227 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 6.0 in stage 6.0 (TID 39, localhost, executor driver, partition 6, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 5.0 in stage 6.0 (TID 38) in 15 ms on localhost (executor driver) (6/10)
21/03/21 21:34:49 INFO Executor: Running task 6.0 in stage 6.0 (TID 39)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms

21/03/21 21:34:49 INFO Executor: Finished task 6.0 in stage 6.0 (TID 39).
1227 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 7.0 in stage 6.0 (TID
40, localhost, executor driver, partition 7, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO Executor: Running task 7.0 in stage 6.0 (TID 40)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:34:49 INFO TaskSetManager: Finished task 6.0 in stage 6.0 (TID
39) in 23 ms on localhost (executor driver) (7/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 1 ms
21/03/21 21:34:49 INFO BlockManagerInfo: Removed broadcast_6_piece0 on ip-
172-20-5-242.us-east-2.compute.internal:44177 in memory (size: 4.2 KB,
free: 424.4 MB)
21/03/21 21:34:49 INFO Executor: Finished task 7.0 in stage 6.0 (TID 40).
1227 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 8.0 in stage 6.0 (TID
41, localhost, executor driver, partition 8, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO Executor: Running task 8.0 in stage 6.0 (TID 41)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 7.0 in stage 6.0 (TID
40) in 11 ms on localhost (executor driver) (8/10)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 2 non-empty
blocks including 2 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 1 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 2 non-empty
blocks including 2 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:34:49 INFO Executor: Finished task 8.0 in stage 6.0 (TID 41).
1356 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Starting task 9.0 in stage 6.0 (TID
42, localhost, executor driver, partition 9, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:34:49 INFO TaskSetManager: Finished task 8.0 in stage 6.0 (TID
41) in 25 ms on localhost (executor driver) (9/10)
21/03/21 21:34:49 INFO Executor: Running task 9.0 in stage 6.0 (TID 42)
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 1 ms
21/03/21 21:34:49 INFO Executor: Finished task 9.0 in stage 6.0 (TID 42).
1356 bytes result sent to driver
21/03/21 21:34:49 INFO TaskSetManager: Finished task 9.0 in stage 6.0 (TID
42) in 25 ms on localhost (executor driver) (10/10)
21/03/21 21:34:49 INFO TaskSchedulerImpl: Removed TaskSet 6.0, whose tasks
have all completed, from pool

21/03/21 21:34:49 INFO DAGScheduler: ShuffleMapStage 6 (repartition at Driver.scala:28) finished in 0.242 s
21/03/21 21:34:49 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:34:49 INFO DAGScheduler: running: Set()
21/03/21 21:34:49 INFO DAGScheduler: waiting: Set(ResultStage 7)
21/03/21 21:34:49 INFO DAGScheduler: failed: Set()
21/03/21 21:34:49 INFO DAGScheduler: Submitting ResultStage 7 (MapPartitionsRDD[23] at saveAsTextFile at Driver.scala:28), which has no missing parents
21/03/21 21:34:49 INFO MemoryStore: Block broadcast_8 stored as values in memory (estimated size 73.6 KB, free 424.1 MB)
21/03/21 21:34:49 INFO MemoryStore: Block broadcast_8_piece0 stored as bytes in memory (estimated size 27.8 KB, free 424.1 MB)
21/03/21 21:34:49 INFO BlockManagerInfo: Added broadcast_8_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:44177 (size: 27.8 KB, free: 424.4 MB)
21/03/21 21:34:49 INFO SparkContext: Created broadcast 8 from broadcast at DAGScheduler.scala:1201
21/03/21 21:34:49 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 7 (MapPartitionsRDD[23] at saveAsTextFile at Driver.scala:28) (first 15 tasks are for partitions Vector(0))
21/03/21 21:34:49 INFO TaskSchedulerImpl: Adding task set 7.0 with 1 tasks
21/03/21 21:34:49 INFO TaskSetManager: Starting task 0.0 in stage 7.0 (TID 43, localhost, executor driver, partition 0, ANY, 7938 bytes)
21/03/21 21:34:49 INFO Executor: Running task 0.0 in stage 7.0 (TID 43)
21/03/21 21:34:49 INFO HadoopMapRedCommitProtocol: Using output committer class org.apache.hadoop.mapred.DirectFileOutputCommitter
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Getting 2 non-empty blocks including 2 local blocks and 0 remote blocks
21/03/21 21:34:49 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 4 ms
21/03/21 21:34:49 INFO MultipartUploadOutputStream: close closed:false s3://801198913bucket/FPGrowthOutput/part-00000
21/03/21 21:34:50 INFO SparkHadoopMapRedUtil: No need to commit output of task because needsTaskCommit=false: attempt_20191027213448_0023_m_000000_0
21/03/21 21:34:50 INFO Executor: Finished task 0.0 in stage 7.0 (TID 43). 1373 bytes result sent to driver
21/03/21 21:34:50 INFO TaskSetManager: Finished task 0.0 in stage 7.0 (TID 43) in 314 ms on localhost (executor driver) (1/1)
21/03/21 21:34:50 INFO TaskSchedulerImpl: Removed TaskSet 7.0, whose tasks have all completed, from pool
21/03/21 21:34:50 INFO DAGScheduler: ResultStage 7 (runJob at SparkHadoopWriter.scala:78) finished in 0.366 s
21/03/21 21:34:50 INFO DAGScheduler: Job 2 finished: runJob at SparkHadoopWriter.scala:78, took 1.222073 s
21/03/21 21:34:50 INFO FileOutputCommitter: File Output Committer Algorithm version is 2
21/03/21 21:34:50 INFO FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: true
21/03/21 21:34:50 INFO DirectFileOutputCommitter: Direct Write: ENABLED
21/03/21 21:34:50 INFO DirectFileOutputCommitter: Nothing to clean up since no temporary files were written.

21/03/21 21:34:50 INFO MultipartUploadOutputStream: close closed:false
s3://801198913bucket/FPGrowthOutput/_SUCCESS
21/03/21 21:34:50 INFO SparkHadoopWriter: Job job_20191027213448_0023
committed.
21/03/21 21:34:50 INFO SparkContext: Invoking stop() from shutdown hook
21/03/21 21:34:50 INFO SparkUI: Stopped Spark web UI at http://ip-172-31-
30-242.us-east-2.compute.internal:4040
21/03/21 21:34:50 INFO ContextCleaner: Cleaned accumulator 93
21/03/21 21:34:50 INFO MapOutputTrackerMasterEndpoint:
MapOutputTrackerMasterEndpoint stopped!
21/03/21 21:34:50 INFO MemoryStore: MemoryStore cleared
21/03/21 21:34:50 INFO BlockManager: BlockManager stopped
21/03/21 21:34:50 INFO BlockManagerMaster: BlockManagerMaster stopped
21/03/21 21:34:50 INFO
OutputCommitCoordinator\$OutputCommitCoordinatorEndpoint:
OutputCommitCoordinator stopped!
21/03/21 21:34:50 INFO SparkContext: Successfully stopped SparkContext
21/03/21 21:34:50 INFO ShutdownHookManager: Shutdown hook called
21/03/21 21:34:50 INFO ShutdownHookManager: Deleting directory
/mnt/tmp/spark-164e8b9d-53c8-4bed-b20b-c5b1e05390c7
21/03/21 21:34:50 INFO ShutdownHookManager: Deleting directory
/mnt/tmp/spark-be2b2756-6ddf-4cae-95ba-49e068570512
[hadoop@ip-172-20-5-242 ~]\$ spark-submit --class org.FP.Driver
./FPGrowth.jar s3://801198913bucket/SampleData.txt 0.5 0.5
s3://801198913bucket/FPGrowthOutput1
21/03/21 21:45:25 WARN SparkConf: The configuration key
'spark.kryoserializer.buffer.max.mb' has been deprecated as of Spark 1.4
and may be removed in the future. Please use the new key
'spark.kryoserializer.buffer.max' instead.
21/03/21 21:45:25 INFO SparkContext: Running Spark version 2.4.3
21/03/21 21:45:25 INFO SparkContext: Submitted application: Association
Action Rules Spark
21/03/21 21:45:25 INFO SecurityManager: Changing view acls to: hadoop
21/03/21 21:45:25 INFO SecurityManager: Changing modify acls to: hadoop
21/03/21 21:45:25 INFO SecurityManager: Changing view acls groups to:
21/03/21 21:45:25 INFO SecurityManager: Changing modify acls groups to:
21/03/21 21:45:25 INFO SecurityManager: SecurityManager: authentication
disabled; ui acls disabled; users with view permissions: Set(hadoop);
groups with view permissions: Set(); users with modify permissions:
Set(hadoop); groups with modify permissions: Set()
21/03/21 21:45:25 INFO Utils: Successfully started service 'sparkDriver'
on port 35919.
21/03/21 21:45:25 INFO SparkEnv: Registering MapOutputTracker
21/03/21 21:45:25 INFO SparkEnv: Registering BlockManagerMaster
21/03/21 21:45:25 INFO BlockManagerMasterEndpoint: Using
org.apache.spark.storage.DefaultTopologyMapper for getting topology
information
21/03/21 21:45:25 INFO BlockManagerMasterEndpoint:
BlockManagerMasterEndpoint up
21/03/21 21:45:25 INFO DiskBlockManager: Created local directory at
/mnt/tmp/blockmgr-a9c492a9-6408-4a35-804a-93542fc6bd0d
21/03/21 21:45:25 INFO MemoryStore: MemoryStore started with capacity
424.4 MB
21/03/21 21:45:26 INFO SparkEnv: Registering OutputCommitCoordinator

21/03/21 21:45:26 INFO Utils: Successfully started service 'SparkUI' on port 4040.
21/03/21 21:45:26 INFO SparkUI: Bound SparkUI to 0.0.0.0, and started at http://ip-172-20-5-242.us-east-2.compute.internal:4040
21/03/21 21:45:26 INFO SparkContext: Added JAR file:/home/hadoop/./FPGrowth.jar at spark://ip-172-20-5-242.us-east-2.compute.internal:35919/jars/FPGrowth.jar with timestamp 1572212726470
21/03/21 21:45:26 INFO Executor: Starting executor ID driver on host localhost
21/03/21 21:45:26 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 46427.
21/03/21 21:45:26 INFO NettyBlockTransferService: Server created on ip-172-20-5-242.us-east-2.compute.internal:46427
21/03/21 21:45:26 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy
21/03/21 21:45:26 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None)
21/03/21 21:45:26 INFO BlockManagerMasterEndpoint: Registering block manager ip-172-20-5-242.us-east-2.compute.internal:46427 with 424.4 MB RAM, BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None)
21/03/21 21:45:26 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None)
21/03/21 21:45:26 INFO BlockManager: external shuffle service port = 7337
21/03/21 21:45:26 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None)
21/03/21 21:45:26 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:26 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering with master
21/03/21 21:45:26 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None)
21/03/21 21:45:26 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None)
21/03/21 21:45:26 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:26 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:26 INFO BlockManager: BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering with master
21/03/21 21:45:26 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None)
21/03/21 21:45:26 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427, None)
21/03/21 21:45:26 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:27 INFO Executor: Told to re-register on heartbeat

[illegible]

[illegible]


```

21/03/21 21:45:27 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)

```

21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO Executor: Told to re-register on heartbeat
21/03/21 21:45:28 INFO BlockManager: BlockManager BlockManagerId(driver,
ip-172-20-5-242.us-east-2.compute.internal, 46427, None) re-registering
with master

```
21/03/21 21:45:28 INFO BlockManagerMaster: Registering BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManagerMaster: Registered BlockManager
BlockManagerId(driver, ip-172-20-5-242.us-east-2.compute.internal, 46427,
None)
21/03/21 21:45:28 INFO BlockManager: Reporting 0 blocks to the master.
21/03/21 21:45:28 INFO EventLoggingListener: Logging events to
hdfs://var/log/spark/apps/local-1572212726558
21/03/21 21:45:29 WARN NettyRpcEnv: Ignored message:
HeartbeatResponse(false)
21/03/21 21:45:29 WARN Executor: Issue communicating with driver in
heartbeater
org.apache.spark.rpc.RpcTimeoutException: Futures timed out after [100
milliseconds]. This timeout is controlled by
spark.executor.heartbeatInterval
    at
org.apache.spark.rpc.RpcTimeout.org$apache$spark$rpc$RpcTimeout$$createRpc
TimeoutException(RpcTimeout.scala:47)
    at
org.apache.spark.rpc.RpcTimeout$$anonfun$addMessageIfTimeout$1.applyOrElse
(RpcTimeout.scala:62)
    at
org.apache.spark.rpc.RpcTimeout$$anonfun$addMessageIfTimeout$1.applyOrElse
(RpcTimeout.scala:58)
    at
scala.runtime.AbstractPartialFunction.apply(AbstractPartialFunction.scala:
36)
    at
org.apache.spark.rpc.RpcTimeout.awaitResult(RpcTimeout.scala:76)
    at
org.apache.spark.rpc.RpcEndpointRef.askSync(RpcEndpointRef.scala:92)
    at
org.apache.spark.executor.Executor.org$apache$spark$executor$Executor$$rep
ortHeartBeat(Executor.scala:841)
    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply$mcV$sp(Exe
cutor.scala:870)
    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply(Executor.s
cala:870)
    at
org.apache.spark.executor.Executor$$anon$2$$anonfun$run$1.apply(Executor.s
cala:870)
    at
org.apache.spark.util.Utils$.logUncaughtExceptions(Utils.scala:1945)
    at
org.apache.spark.executor.Executor$$anon$2.run(Executor.scala:870)
    at
java.util.concurrent.Executors$RunnableAdapter.call(Executors.java:511)
    at
java.util.concurrent.FutureTask.runAndReset(FutureTask.java:308)
```

```

    at
java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.access$301(ScheduledThreadPoolExecutor.java:180)
    at
java.util.concurrent.ScheduledThreadPoolExecutor$ScheduledFutureTask.run(ScheduledThreadPoolExecutor.java:294)
    at
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
    at
java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
    at java.lang.Thread.run(Thread.java:748)
Caused by: java.util.concurrent.TimeoutException: Futures timed out after [100 milliseconds]
    at
scala.concurrent.impl.Promise$DefaultPromise.ready(Promise.scala:223)
    at
scala.concurrent.impl.Promise$DefaultPromise.result(Promise.scala:227)
    at
org.apache.spark.util.ThreadUtils$.awaitResult(ThreadUtils.scala:220)
    at
org.apache.spark.rpc.RpcTimeout.awaitResult(RpcTimeout.scala:75)
    ... 14 more
21/03/21 21:45:29 INFO MemoryStore: Block broadcast_0 stored as values in memory (estimated size 237.8 KB, free 424.2 MB)
21/03/21 21:45:29 INFO MemoryStore: Block broadcast_0_piece0 stored as bytes in memory (estimated size 24.2 KB, free 424.2 MB)
21/03/21 21:45:29 INFO BlockManagerInfo: Added broadcast_0_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 24.2 KB, free: 424.4 MB)
21/03/21 21:45:29 INFO SparkContext: Created broadcast 0 from textFile at Driver.scala:18
21/03/21 21:45:30 WARN FPGrowth: Input data is not cached.
21/03/21 21:45:30 INFO GPLNativeCodeLoader: Loaded native gpl library
21/03/21 21:45:30 INFO LzoCodec: Successfully loaded & initialized native-lzo library [hadoop-lzo rev 7e6c862e89bc8db32c064454a55af74ddff73bae]
21/03/21 21:45:32 INFO FileInputFormat: Total input files to process : 1
21/03/21 21:45:32 INFO SparkContext: Starting job: count at FPGrowth.scala:217
21/03/21 21:45:32 INFO DAGScheduler: Got job 0 (count at FPGrowth.scala:217) with 1 output partitions
21/03/21 21:45:32 INFO DAGScheduler: Final stage: ResultStage 0 (count at FPGrowth.scala:217)
21/03/21 21:45:32 INFO DAGScheduler: Parents of final stage: List()
21/03/21 21:45:32 INFO DAGScheduler: Missing parents: List()
21/03/21 21:45:32 INFO DAGScheduler: Submitting ResultStage 0 (MapPartitionsRDD[2] at map at Driver.scala:20), which has no missing parents
21/03/21 21:45:32 INFO MemoryStore: Block broadcast_1 stored as values in memory (estimated size 3.4 KB, free 424.2 MB)
21/03/21 21:45:32 INFO MemoryStore: Block broadcast_1_piece0 stored as bytes in memory (estimated size 2.1 KB, free 424.2 MB)

```

21/03/21 21:45:32 INFO BlockManagerInfo: Added broadcast_1_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 2.1 KB, free: 424.4 MB)

21/03/21 21:45:32 INFO SparkContext: Created broadcast 1 from broadcast at DAGScheduler.scala:1201

21/03/21 21:45:32 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 0 (MapPartitionsRDD[2] at map at Driver.scala:20) (first 15 tasks are for partitions Vector(0))

21/03/21 21:45:32 INFO TaskSchedulerImpl: Adding task set 0.0 with 1 tasks

21/03/21 21:45:32 INFO TaskSetManager: Starting task 0.0 in stage 0.0 (TID 0, localhost, executor driver, partition 0, ANY, 7895 bytes)

21/03/21 21:45:32 INFO Executor: Running task 0.0 in stage 0.0 (TID 0)

21/03/21 21:45:32 INFO Executor: Fetching spark://ip-172-20-5-242.us-east-2.compute.internal:35919/jars/FPGrowth.jar with timestamp 1572212726470

21/03/21 21:45:32 INFO TransportClientFactory: Successfully created connection to ip-172-20-5-242.us-east-2.compute.internal/172.31.30.242:35919 after 71 ms (0 ms spent in bootstraps)

21/03/21 21:45:32 INFO Utils: Fetching spark://ip-172-20-5-242.us-east-2.compute.internal:35919/jars/FPGrowth.jar to /mnt/tmp/spark-ab4282d4-3721-48e6-9762-35d60b619cf6/userFiles-84d9719d-3e55-432d-8833-77f3bc6f786b/fetchFileTemp1056233876874591175.tmp

21/03/21 21:45:33 INFO Executor: Adding file:/mnt/tmp/spark-ab4282d4-3721-48e6-9762-35d60b619cf6/userFiles-84d9719d-3e55-432d-8833-77f3bc6f786b/FPGrowth.jar to class loader

21/03/21 21:45:33 INFO HadoopRDD: Input split: s3://801198913bucket/SampleData.txt:0+63

21/03/21 21:45:33 INFO S3NativeFileSystem: Opening 's3://801198913bucket/SampleData.txt' for reading

21/03/21 21:45:33 INFO Executor: Finished task 0.0 in stage 0.0 (TID 0). 918 bytes result sent to driver

21/03/21 21:45:33 INFO TaskSetManager: Finished task 0.0 in stage 0.0 (TID 0) in 619 ms on localhost (executor driver) (1/1)

21/03/21 21:45:33 INFO DAGScheduler: ResultStage 0 (count at FPGrowth.scala:217) finished in 0.828 s

21/03/21 21:45:33 INFO TaskSchedulerImpl: Removed TaskSet 0.0, whose tasks have all completed, from pool

21/03/21 21:45:33 INFO DAGScheduler: Job 0 finished: count at FPGrowth.scala:217, took 1.048029 s

21/03/21 21:45:33 INFO BlockManagerInfo: Removed broadcast_1_piece0 on ip-172-20-5-242.us-east-2.compute.internal:46427 in memory (size: 2.1 KB, free: 424.4 MB)

21/03/21 21:45:33 INFO SparkContext: Starting job: collect at FPGrowth.scala:257

21/03/21 21:45:34 INFO DAGScheduler: Registering RDD 4 (map at FPGrowth.scala:254)

21/03/21 21:45:34 INFO DAGScheduler: Got job 1 (collect at FPGrowth.scala:257) with 10 output partitions

21/03/21 21:45:34 INFO DAGScheduler: Final stage: ResultStage 2 (collect at FPGrowth.scala:257)

21/03/21 21:45:34 INFO DAGScheduler: Parents of final stage: List(ShuffleMapStage 1)

21/03/21 21:45:34 INFO DAGScheduler: Missing parents: List(ShuffleMapStage 1)
21/03/21 21:45:34 INFO DAGScheduler: Submitting ShuffleMapStage 1 (MapPartitionsRDD[4] at map at FPGrowth.scala:254), which has no missing parents
21/03/21 21:45:34 INFO MemoryStore: Block broadcast_2 stored as values in memory (estimated size 5.1 KB, free 424.2 MB)
21/03/21 21:45:34 INFO MemoryStore: Block broadcast_2_piece0 stored as bytes in memory (estimated size 2.9 KB, free 424.2 MB)
21/03/21 21:45:34 INFO BlockManagerInfo: Added broadcast_2_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 2.9 KB, free: 424.4 MB)
21/03/21 21:45:34 INFO SparkContext: Created broadcast 2 from broadcast at DAGScheduler.scala:1201
21/03/21 21:45:34 INFO DAGScheduler: Submitting 1 missing tasks from ShuffleMapStage 1 (MapPartitionsRDD[4] at map at FPGrowth.scala:254) (first 15 tasks are for partitions Vector(0))
21/03/21 21:45:34 INFO TaskSchedulerImpl: Adding task set 1.0 with 1 tasks
21/03/21 21:45:34 INFO TaskSetManager: Starting task 0.0 in stage 1.0 (TID 1, localhost, executor driver, partition 0, ANY, 7884 bytes)
21/03/21 21:45:34 INFO Executor: Running task 0.0 in stage 1.0 (TID 1)
21/03/21 21:45:34 INFO HadoopRDD: Input split: s3://801198913bucket/SampleData.txt:0+63
21/03/21 21:45:34 INFO S3NativeFileSystem: Opening 's3://801198913bucket/SampleData.txt' for reading
21/03/21 21:45:34 INFO Executor: Finished task 0.0 in stage 1.0 (TID 1). 1193 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Finished task 0.0 in stage 1.0 (TID 1) in 212 ms on localhost (executor driver) (1/1)
21/03/21 21:45:34 INFO TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks have all completed, from pool
21/03/21 21:45:34 INFO DAGScheduler: ShuffleMapStage 1 (map at FPGrowth.scala:254) finished in 0.255 s
21/03/21 21:45:34 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:45:34 INFO DAGScheduler: running: Set()
21/03/21 21:45:34 INFO DAGScheduler: waiting: Set(ResultStage 2)
21/03/21 21:45:34 INFO DAGScheduler: failed: Set()
21/03/21 21:45:34 INFO DAGScheduler: Submitting ResultStage 2 (MapPartitionsRDD[6] at filter at FPGrowth.scala:256), which has no missing parents
21/03/21 21:45:34 INFO MemoryStore: Block broadcast_3 stored as values in memory (estimated size 3.8 KB, free 424.2 MB)
21/03/21 21:45:34 INFO MemoryStore: Block broadcast_3_piece0 stored as bytes in memory (estimated size 2.3 KB, free 424.2 MB)
21/03/21 21:45:34 INFO BlockManagerInfo: Added broadcast_3_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 2.3 KB, free: 424.4 MB)
21/03/21 21:45:34 INFO SparkContext: Created broadcast 3 from broadcast at DAGScheduler.scala:1201
21/03/21 21:45:34 INFO DAGScheduler: Submitting 10 missing tasks from ResultStage 2 (MapPartitionsRDD[6] at filter at FPGrowth.scala:256) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9))
21/03/21 21:45:34 INFO TaskSchedulerImpl: Adding task set 2.0 with 10 tasks

21/03/21 21:45:34 INFO TaskSetManager: Starting task 3.0 in stage 2.0 (TID 2, localhost, executor driver, partition 3, PROCESS_LOCAL, 7662 bytes)
21/03/21 21:45:34 INFO Executor: Running task 3.0 in stage 2.0 (TID 2)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 12 ms
21/03/21 21:45:34 INFO Executor: Finished task 3.0 in stage 2.0 (TID 2). 1134 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Starting task 4.0 in stage 2.0 (TID 3, localhost, executor driver, partition 4, PROCESS_LOCAL, 7662 bytes)
21/03/21 21:45:34 INFO Executor: Running task 4.0 in stage 2.0 (TID 3)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:34 INFO Executor: Finished task 4.0 in stage 2.0 (TID 3). 1134 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Starting task 5.0 in stage 2.0 (TID 4, localhost, executor driver, partition 5, PROCESS_LOCAL, 7662 bytes)
21/03/21 21:45:34 INFO TaskSetManager: Finished task 3.0 in stage 2.0 (TID 2) in 76 ms on localhost (executor driver) (1/10)
21/03/21 21:45:34 INFO Executor: Running task 5.0 in stage 2.0 (TID 4)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:34 INFO TaskSetManager: Finished task 4.0 in stage 2.0 (TID 3) in 23 ms on localhost (executor driver) (2/10)
21/03/21 21:45:34 INFO Executor: Finished task 5.0 in stage 2.0 (TID 4). 1134 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Starting task 0.0 in stage 2.0 (TID 5, localhost, executor driver, partition 0, ANY, 7662 bytes)
21/03/21 21:45:34 INFO TaskSetManager: Finished task 5.0 in stage 2.0 (TID 4) in 22 ms on localhost (executor driver) (3/10)
21/03/21 21:45:34 INFO Executor: Running task 0.0 in stage 2.0 (TID 5)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 2 ms
21/03/21 21:45:34 INFO Executor: Finished task 0.0 in stage 2.0 (TID 5). 1134 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Starting task 1.0 in stage 2.0 (TID 6, localhost, executor driver, partition 1, ANY, 7662 bytes)
21/03/21 21:45:34 INFO TaskSetManager: Finished task 0.0 in stage 2.0 (TID 5) in 45 ms on localhost (executor driver) (4/10)
21/03/21 21:45:34 INFO Executor: Running task 1.0 in stage 2.0 (TID 6)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:34 INFO Executor: Finished task 1.0 in stage 2.0 (TID 6). 1134 bytes result sent to driver

21/03/21 21:45:34 INFO TaskSetManager: Starting task 2.0 in stage 2.0 (TID 7, localhost, executor driver, partition 2, ANY, 7662 bytes)
21/03/21 21:45:34 INFO TaskSetManager: Finished task 1.0 in stage 2.0 (TID 6) in 16 ms on localhost (executor driver) (5/10)
21/03/21 21:45:34 INFO Executor: Running task 2.0 in stage 2.0 (TID 7)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:34 INFO Executor: Finished task 2.0 in stage 2.0 (TID 7). 1134 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Starting task 6.0 in stage 2.0 (TID 8, localhost, executor driver, partition 6, ANY, 7662 bytes)
21/03/21 21:45:34 INFO TaskSetManager: Finished task 2.0 in stage 2.0 (TID 7) in 16 ms on localhost (executor driver) (6/10)
21/03/21 21:45:34 INFO Executor: Running task 6.0 in stage 2.0 (TID 8)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:34 INFO Executor: Finished task 6.0 in stage 2.0 (TID 8). 1282 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Starting task 7.0 in stage 2.0 (TID 9, localhost, executor driver, partition 7, ANY, 7662 bytes)
21/03/21 21:45:34 INFO TaskSetManager: Finished task 6.0 in stage 2.0 (TID 8) in 24 ms on localhost (executor driver) (7/10)
21/03/21 21:45:34 INFO Executor: Running task 7.0 in stage 2.0 (TID 9)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:34 INFO Executor: Finished task 7.0 in stage 2.0 (TID 9). 1282 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Starting task 8.0 in stage 2.0 (TID 10, localhost, executor driver, partition 8, ANY, 7662 bytes)
21/03/21 21:45:34 INFO TaskSetManager: Finished task 7.0 in stage 2.0 (TID 9) in 18 ms on localhost (executor driver) (8/10)
21/03/21 21:45:34 INFO Executor: Running task 8.0 in stage 2.0 (TID 10)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:34 INFO Executor: Finished task 8.0 in stage 2.0 (TID 10). 1282 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Starting task 9.0 in stage 2.0 (TID 11, localhost, executor driver, partition 9, ANY, 7662 bytes)
21/03/21 21:45:34 INFO TaskSetManager: Finished task 8.0 in stage 2.0 (TID 10) in 20 ms on localhost (executor driver) (9/10)
21/03/21 21:45:34 INFO Executor: Running task 9.0 in stage 2.0 (TID 11)
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:34 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms

21/03/21 21:45:34 INFO Executor: Finished task 9.0 in stage 2.0 (TID 11).
1282 bytes result sent to driver
21/03/21 21:45:34 INFO TaskSetManager: Finished task 9.0 in stage 2.0 (TID
11) in 19 ms on localhost (executor driver) (10/10)
21/03/21 21:45:34 INFO TaskSchedulerImpl: Removed TaskSet 2.0, whose tasks
have all completed, from pool
21/03/21 21:45:34 INFO DAGScheduler: ResultStage 2 (collect at
FPGrowth.scala:257) finished in 0.255 s
21/03/21 21:45:34 INFO DAGScheduler: Job 1 finished: collect at
FPGrowth.scala:257, took 1.069130 s
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 36
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 70
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 47
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 25
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 68
21/03/21 21:45:34 INFO BlockManagerInfo: Removed broadcast_2_piece0 on ip-
172-20-5-242.us-east-2.compute.internal:46427 in memory (size: 2.9 KB,
free: 424.4 MB)
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 33
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 69
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 64
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 61
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 42
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 28
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 53
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 48
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 73
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 44
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 49
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 52
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 62
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 67
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 38
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 63
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 74
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 41
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 40
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 31
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 27
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 65
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 51
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 66
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 55
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 57
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 45
21/03/21 21:45:34 INFO BlockManagerInfo: Removed broadcast_3_piece0 on ip-
172-20-5-242.us-east-2.compute.internal:46427 in memory (size: 2.3 KB,
free: 424.4 MB)
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 72
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 39
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 58
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 30
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 56
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 60

21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 35
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 34
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 29
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 46
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 71
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 54
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 32
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 50
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 26
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 59
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 37
21/03/21 21:45:34 INFO ContextCleaner: Cleaned shuffle 0
21/03/21 21:45:34 INFO ContextCleaner: Cleaned accumulator 43
21/03/21 21:45:35 INFO deprecation: mapred.output.dir is deprecated.
Instead, use mapreduce.output.fileoutputformat.outputdir
21/03/21 21:45:35 INFO HadoopMapRedCommitProtocol: Using output committer
class org.apache.hadoop.mapred.DirectFileOutputCommitter
21/03/21 21:45:35 INFO DirectFileOutputCommitter: Nothing to setup since
the outputs are written directly.
21/03/21 21:45:35 INFO SparkContext: Starting job: runJob at
SparkHadoopWriter.scala:78
21/03/21 21:45:35 INFO DAGScheduler: Registering RDD 7 (flatMap at
FPGrowth.scala:275)
21/03/21 21:45:35 INFO DAGScheduler: Registering RDD 11 (flatMap at
AssociationRules.scala:78)
21/03/21 21:45:35 INFO DAGScheduler: Registering RDD 12 (map at
AssociationRules.scala:90)
21/03/21 21:45:35 INFO DAGScheduler: Registering RDD 19 (repartition at
Driver.scala:28)
21/03/21 21:45:35 INFO DAGScheduler: Got job 2 (runJob at
SparkHadoopWriter.scala:78) with 1 output partitions
21/03/21 21:45:35 INFO DAGScheduler: Final stage: ResultStage 7 (runJob at
SparkHadoopWriter.scala:78)
21/03/21 21:45:35 INFO DAGScheduler: Parents of final stage:
List(ShuffleMapStage 6)
21/03/21 21:45:35 INFO DAGScheduler: Missing parents: List(ShuffleMapStage
6)
21/03/21 21:45:35 INFO DAGScheduler: Submitting ShuffleMapStage 3
(MapPartitionsRDD[7] at flatMap at FPGrowth.scala:275), which has no
missing parents
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_4 stored as values in
memory (estimated size 7.1 KB, free 424.2 MB)
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_4_piece0 stored as
bytes in memory (estimated size 3.7 KB, free 424.2 MB)
21/03/21 21:45:35 INFO BlockManagerInfo: Added broadcast_4_piece0 in
memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 3.7 KB,
free: 424.4 MB)
21/03/21 21:45:35 INFO SparkContext: Created broadcast 4 from broadcast at
DAGScheduler.scala:1201
21/03/21 21:45:35 INFO DAGScheduler: Submitting 1 missing tasks from
ShuffleMapStage 3 (MapPartitionsRDD[7] at flatMap at FPGrowth.scala:275)
(first 15 tasks are for partitions Vector(0))
21/03/21 21:45:35 INFO TaskSchedulerImpl: Adding task set 3.0 with 1 tasks

21/03/21 21:45:35 INFO TaskSetManager: Starting task 0.0 in stage 3.0 (TID 12, localhost, executor driver, partition 0, ANY, 7884 bytes)
21/03/21 21:45:35 INFO Executor: Running task 0.0 in stage 3.0 (TID 12)
21/03/21 21:45:35 INFO HadoopRDD: Input split:
s3://801198913bucket/SampleData.txt:0+63
21/03/21 21:45:35 INFO S3NativeFileSystem: Opening
's3://801198913bucket/SampleData.txt' for reading
21/03/21 21:45:35 INFO Executor: Finished task 0.0 in stage 3.0 (TID 12).
1150 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Finished task 0.0 in stage 3.0 (TID 12) in 62 ms on localhost (executor driver) (1/1)
21/03/21 21:45:35 INFO TaskSchedulerImpl: Removed TaskSet 3.0, whose tasks have all completed, from pool
21/03/21 21:45:35 INFO DAGScheduler: ShuffleMapStage 3 (flatMap at FPGrowth.scala:275) finished in 0.088 s
21/03/21 21:45:35 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:45:35 INFO DAGScheduler: running: Set()
21/03/21 21:45:35 INFO DAGScheduler: waiting: Set(ShuffleMapStage 5, ShuffleMapStage 6, ResultStage 7, ShuffleMapStage 4)
21/03/21 21:45:35 INFO DAGScheduler: failed: Set()
21/03/21 21:45:35 INFO DAGScheduler: Submitting ShuffleMapStage 5 (MapPartitionsRDD[12] at map at AssociationRules.scala:90), which has no missing parents
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_5 stored as values in memory (estimated size 8.5 KB, free 424.2 MB)
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_5_piece0 stored as bytes in memory (estimated size 4.3 KB, free 424.1 MB)
21/03/21 21:45:35 INFO BlockManagerInfo: Added broadcast_5_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 4.3 KB, free: 424.4 MB)
21/03/21 21:45:35 INFO SparkContext: Created broadcast 5 from broadcast at DAGScheduler.scala:1201
21/03/21 21:45:35 INFO DAGScheduler: Submitting 10 missing tasks from ShuffleMapStage 5 (MapPartitionsRDD[12] at map at AssociationRules.scala:90) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9))
21/03/21 21:45:35 INFO TaskSchedulerImpl: Adding task set 5.0 with 10 tasks
21/03/21 21:45:35 INFO TaskSetManager: Starting task 4.0 in stage 5.0 (TID 13, localhost, executor driver, partition 4, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO Executor: Running task 4.0 in stage 5.0 (TID 13)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO DAGScheduler: Submitting ShuffleMapStage 4 (MapPartitionsRDD[11] at flatMap at AssociationRules.scala:78), which has no missing parents
21/03/21 21:45:35 INFO Executor: Finished task 4.0 in stage 5.0 (TID 13).
1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 5.0 in stage 5.0 (TID 14, localhost, executor driver, partition 5, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO Executor: Running task 5.0 in stage 5.0 (TID 14)

21/03/21 21:45:35 INFO TaskSetManager: Finished task 4.0 in stage 5.0 (TID 13) in 13 ms on localhost (executor driver) (1/10)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 5.0 in stage 5.0 (TID 14). 1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 6.0 in stage 5.0 (TID 15, localhost, executor driver, partition 6, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 5.0 in stage 5.0 (TID 14) in 6 ms on localhost (executor driver) (2/10)
21/03/21 21:45:35 INFO Executor: Running task 6.0 in stage 5.0 (TID 15)
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_6 stored as values in memory (estimated size 8.5 KB, free 424.1 MB)
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_6_piece0 stored as bytes in memory (estimated size 4.3 KB, free 424.1 MB)
21/03/21 21:45:35 INFO BlockManagerInfo: Added broadcast_6_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 4.3 KB, free: 424.4 MB)
21/03/21 21:45:35 INFO SparkContext: Created broadcast 6 from broadcast at DAGScheduler.scala:1201
21/03/21 21:45:35 INFO DAGScheduler: Submitting 10 missing tasks from ShuffleMapStage 4 (MapPartitionsRDD[11] at flatMap at AssociationRules.scala:78) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9))
21/03/21 21:45:35 INFO TaskSchedulerImpl: Adding task set 4.0 with 10 tasks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 8 ms
21/03/21 21:45:35 INFO Executor: Finished task 6.0 in stage 5.0 (TID 15). 1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 4.0 in stage 4.0 (TID 16, localhost, executor driver, partition 4, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 6.0 in stage 5.0 (TID 15) in 26 ms on localhost (executor driver) (3/10)
21/03/21 21:45:35 INFO Executor: Running task 4.0 in stage 4.0 (TID 16)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 4.0 in stage 4.0 (TID 16). 1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 5.0 in stage 4.0 (TID 17, localhost, executor driver, partition 5, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 4.0 in stage 4.0 (TID 16) in 11 ms on localhost (executor driver) (1/10)
21/03/21 21:45:35 INFO Executor: Running task 5.0 in stage 4.0 (TID 17)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms

21/03/21 21:45:35 INFO Executor: Finished task 5.0 in stage 4.0 (TID 17).
1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 6.0 in stage 4.0 (TID
18, localhost, executor driver, partition 6, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO Executor: Running task 6.0 in stage 4.0 (TID 18)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 5.0 in stage 4.0 (TID
17) in 25 ms on localhost (executor driver) (2/10)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 6.0 in stage 4.0 (TID 18).
1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 7.0 in stage 4.0 (TID
19, localhost, executor driver, partition 7, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 6.0 in stage 4.0 (TID
18) in 8 ms on localhost (executor driver) (3/10)
21/03/21 21:45:35 INFO Executor: Running task 7.0 in stage 4.0 (TID 19)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 7.0 in stage 4.0 (TID 19).
1279 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 8.0 in stage 4.0 (TID
20, localhost, executor driver, partition 8, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO Executor: Running task 8.0 in stage 4.0 (TID 20)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 7.0 in stage 4.0 (TID
19) in 12 ms on localhost (executor driver) (4/10)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 8.0 in stage 4.0 (TID 20).
1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 9.0 in stage 4.0 (TID
21, localhost, executor driver, partition 9, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 8.0 in stage 4.0 (TID
20) in 8 ms on localhost (executor driver) (5/10)
21/03/21 21:45:35 INFO Executor: Running task 9.0 in stage 4.0 (TID 21)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty
blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote
fetches in 7 ms
21/03/21 21:45:35 INFO Executor: Finished task 9.0 in stage 4.0 (TID 21).
1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 0.0 in stage 4.0 (TID
22, localhost, executor driver, partition 0, ANY, 7651 bytes)
21/03/21 21:45:35 INFO Executor: Running task 0.0 in stage 4.0 (TID 22)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 9.0 in stage 4.0 (TID
21) in 15 ms on localhost (executor driver) (6/10)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty
blocks including 1 local blocks and 0 remote blocks

21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO Executor: Finished task 0.0 in stage 4.0 (TID 22). 1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 1.0 in stage 4.0 (TID 23, localhost, executor driver, partition 1, ANY, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 0.0 in stage 4.0 (TID 22) in 23 ms on localhost (executor driver) (7/10)
21/03/21 21:45:35 INFO Executor: Running task 1.0 in stage 4.0 (TID 23)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO Executor: Finished task 1.0 in stage 4.0 (TID 23). 1365 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 2.0 in stage 4.0 (TID 24, localhost, executor driver, partition 2, ANY, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 1.0 in stage 4.0 (TID 23) in 24 ms on localhost (executor driver) (8/10)
21/03/21 21:45:35 INFO Executor: Running task 2.0 in stage 4.0 (TID 24)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 2.0 in stage 4.0 (TID 24). 1365 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 3.0 in stage 4.0 (TID 25, localhost, executor driver, partition 3, ANY, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 2.0 in stage 4.0 (TID 24) in 21 ms on localhost (executor driver) (9/10)
21/03/21 21:45:35 INFO Executor: Running task 3.0 in stage 4.0 (TID 25)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 3.0 in stage 4.0 (TID 25). 1365 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 7.0 in stage 5.0 (TID 26, localhost, executor driver, partition 7, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 3.0 in stage 4.0 (TID 25) in 14 ms on localhost (executor driver) (10/10)
21/03/21 21:45:35 INFO TaskSchedulerImpl: Removed TaskSet 4.0, whose tasks have all completed, from pool
21/03/21 21:45:35 INFO DAGScheduler: ShuffleMapStage 4 (flatMap at AssociationRules.scala:78) finished in 0.189 s
21/03/21 21:45:35 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:45:35 INFO DAGScheduler: running: Set(ShuffleMapStage 5)
21/03/21 21:45:35 INFO DAGScheduler: waiting: Set(ShuffleMapStage 6, ResultStage 7)
21/03/21 21:45:35 INFO DAGScheduler: failed: Set()
21/03/21 21:45:35 INFO Executor: Running task 7.0 in stage 5.0 (TID 26)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks

21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO Executor: Finished task 7.0 in stage 5.0 (TID 26). 1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 8.0 in stage 5.0 (TID 27, localhost, executor driver, partition 8, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 7.0 in stage 5.0 (TID 26) in 9 ms on localhost (executor driver) (4/10)
21/03/21 21:45:35 INFO Executor: Running task 8.0 in stage 5.0 (TID 27)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 8.0 in stage 5.0 (TID 27). 1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 9.0 in stage 5.0 (TID 28, localhost, executor driver, partition 9, PROCESS_LOCAL, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 8.0 in stage 5.0 (TID 27) in 15 ms on localhost (executor driver) (5/10)
21/03/21 21:45:35 INFO Executor: Running task 9.0 in stage 5.0 (TID 28)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO Executor: Finished task 9.0 in stage 5.0 (TID 28). 1236 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 0.0 in stage 5.0 (TID 29, localhost, executor driver, partition 0, ANY, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 9.0 in stage 5.0 (TID 28) in 12 ms on localhost (executor driver) (6/10)
21/03/21 21:45:35 INFO Executor: Running task 0.0 in stage 5.0 (TID 29)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 0.0 in stage 5.0 (TID 29). 1365 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 1.0 in stage 5.0 (TID 30, localhost, executor driver, partition 1, ANY, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 0.0 in stage 5.0 (TID 29) in 13 ms on localhost (executor driver) (7/10)
21/03/21 21:45:35 INFO Executor: Running task 1.0 in stage 5.0 (TID 30)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 1.0 in stage 5.0 (TID 30). 1365 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 2.0 in stage 5.0 (TID 31, localhost, executor driver, partition 2, ANY, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 1.0 in stage 5.0 (TID 30) in 16 ms on localhost (executor driver) (8/10)
21/03/21 21:45:35 INFO Executor: Running task 2.0 in stage 5.0 (TID 31)

21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 2.0 in stage 5.0 (TID 31). 1365 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 3.0 in stage 5.0 (TID 32, localhost, executor driver, partition 3, ANY, 7651 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 2.0 in stage 5.0 (TID 31) in 18 ms on localhost (executor driver) (9/10)
21/03/21 21:45:35 INFO Executor: Running task 3.0 in stage 5.0 (TID 32)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO Executor: Finished task 3.0 in stage 5.0 (TID 32). 1365 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Finished task 3.0 in stage 5.0 (TID 32) in 15 ms on localhost (executor driver) (10/10)
21/03/21 21:45:35 INFO TaskSchedulerImpl: Removed TaskSet 5.0, whose tasks have all completed, from pool
21/03/21 21:45:35 INFO DAGScheduler: ShuffleMapStage 5 (map at AssociationRules.scala:90) finished in 0.316 s
21/03/21 21:45:35 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:45:35 INFO DAGScheduler: running: Set()
21/03/21 21:45:35 INFO DAGScheduler: waiting: Set(ShuffleMapStage 6, ResultStage 7)
21/03/21 21:45:35 INFO DAGScheduler: failed: Set()
21/03/21 21:45:35 INFO DAGScheduler: Submitting ShuffleMapStage 6 (MapPartitionsRDD[19] at repartition at Driver.scala:28), which has no missing parents
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_7 stored as values in memory (estimated size 5.6 KB, free 424.1 MB)
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_7_piece0 stored as bytes in memory (estimated size 3.0 KB, free 424.1 MB)
21/03/21 21:45:35 INFO BlockManagerInfo: Added broadcast_7_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 3.0 KB, free: 424.4 MB)
21/03/21 21:45:35 INFO SparkContext: Created broadcast 7 from broadcast at DAGScheduler.scala:1201
21/03/21 21:45:35 INFO DAGScheduler: Submitting 10 missing tasks from ShuffleMapStage 6 (MapPartitionsRDD[19] at repartition at Driver.scala:28) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9))
21/03/21 21:45:35 INFO TaskSchedulerImpl: Adding task set 6.0 with 10 tasks
21/03/21 21:45:35 INFO TaskSetManager: Starting task 0.0 in stage 6.0 (TID 33, localhost, executor driver, partition 0, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO Executor: Running task 0.0 in stage 6.0 (TID 33)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks

21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO Executor: Finished task 0.0 in stage 6.0 (TID 33). 1227 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 1.0 in stage 6.0 (TID 34, localhost, executor driver, partition 1, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 0.0 in stage 6.0 (TID 33) in 29 ms on localhost (executor driver) (1/10)
21/03/21 21:45:35 INFO Executor: Running task 1.0 in stage 6.0 (TID 34)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 2 non-empty blocks including 2 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO BlockManagerInfo: Removed broadcast_5_piece0 on ip-172-20-5-242.us-east-2.compute.internal:46427 in memory (size: 4.3 KB, free: 424.4 MB)
21/03/21 21:45:35 INFO Executor: Finished task 1.0 in stage 6.0 (TID 34). 1356 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 2.0 in stage 6.0 (TID 35, localhost, executor driver, partition 2, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO BlockManagerInfo: Removed broadcast_6_piece0 on ip-172-20-5-242.us-east-2.compute.internal:46427 in memory (size: 4.3 KB, free: 424.4 MB)
21/03/21 21:45:35 INFO BlockManagerInfo: Removed broadcast_4_piece0 on ip-172-20-5-242.us-east-2.compute.internal:46427 in memory (size: 3.7 KB, free: 424.4 MB)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 1.0 in stage 6.0 (TID 34) in 55 ms on localhost (executor driver) (2/10)
21/03/21 21:45:35 INFO Executor: Running task 2.0 in stage 6.0 (TID 35)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 3 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 2.0 in stage 6.0 (TID 35). 1227 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 3.0 in stage 6.0 (TID 36, localhost, executor driver, partition 3, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO Executor: Running task 3.0 in stage 6.0 (TID 36)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 2.0 in stage 6.0 (TID 35) in 27 ms on localhost (executor driver) (3/10)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks

21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 3.0 in stage 6.0 (TID 36). 1227 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 4.0 in stage 6.0 (TID 37, localhost, executor driver, partition 4, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 3.0 in stage 6.0 (TID 36) in 8 ms on localhost (executor driver) (4/10)
21/03/21 21:45:35 INFO Executor: Running task 4.0 in stage 6.0 (TID 37)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 4.0 in stage 6.0 (TID 37). 1227 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 5.0 in stage 6.0 (TID 38, localhost, executor driver, partition 5, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 4.0 in stage 6.0 (TID 37) in 10 ms on localhost (executor driver) (5/10)
21/03/21 21:45:35 INFO Executor: Running task 5.0 in stage 6.0 (TID 38)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 5.0 in stage 6.0 (TID 38). 1356 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 6.0 in stage 6.0 (TID 39, localhost, executor driver, partition 6, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 5.0 in stage 6.0 (TID 38) in 23 ms on localhost (executor driver) (6/10)
21/03/21 21:45:35 INFO Executor: Running task 6.0 in stage 6.0 (TID 39)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO Executor: Finished task 6.0 in stage 6.0 (TID 39). 1227 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 7.0 in stage 6.0 (TID 40, localhost, executor driver, partition 7, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO Executor: Running task 7.0 in stage 6.0 (TID 40)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 6.0 in stage 6.0 (TID 39) in 10 ms on localhost (executor driver) (7/10)

21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 7.0 in stage 6.0 (TID 40). 1356 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 8.0 in stage 6.0 (TID 41, localhost, executor driver, partition 8, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 7.0 in stage 6.0 (TID 40) in 16 ms on localhost (executor driver) (8/10)
21/03/21 21:45:35 INFO Executor: Running task 8.0 in stage 6.0 (TID 41)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 0 non-empty blocks including 0 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 3 non-empty blocks including 3 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 8.0 in stage 6.0 (TID 41). 1227 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Starting task 9.0 in stage 6.0 (TID 42, localhost, executor driver, partition 9, PROCESS_LOCAL, 7714 bytes)
21/03/21 21:45:35 INFO Executor: Running task 9.0 in stage 6.0 (TID 42)
21/03/21 21:45:35 INFO TaskSetManager: Finished task 8.0 in stage 6.0 (TID 41) in 16 ms on localhost (executor driver) (9/10)
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 2 non-empty blocks including 2 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 1 ms
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks including 1 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO Executor: Finished task 9.0 in stage 6.0 (TID 42). 1356 bytes result sent to driver
21/03/21 21:45:35 INFO TaskSetManager: Finished task 9.0 in stage 6.0 (TID 42) in 24 ms on localhost (executor driver) (10/10)
21/03/21 21:45:35 INFO TaskSchedulerImpl: Removed TaskSet 6.0, whose tasks have all completed, from pool
21/03/21 21:45:35 INFO DAGScheduler: ShuffleMapStage 6 (repartition at Driver.scala:28) finished in 0.210 s
21/03/21 21:45:35 INFO DAGScheduler: looking for newly runnable stages
21/03/21 21:45:35 INFO DAGScheduler: running: Set()
21/03/21 21:45:35 INFO DAGScheduler: waiting: Set(ResultStage 7)
21/03/21 21:45:35 INFO DAGScheduler: failed: Set()
21/03/21 21:45:35 INFO DAGScheduler: Submitting ResultStage 7 (MapPartitionsRDD[23] at saveAsTextFile at Driver.scala:28), which has no missing parents

21/03/21 21:45:35 INFO MemoryStore: Block broadcast_8 stored as values in memory (estimated size 73.6 KB, free 424.1 MB)
21/03/21 21:45:35 INFO MemoryStore: Block broadcast_8_piece0 stored as bytes in memory (estimated size 27.8 KB, free 424.1 MB)
21/03/21 21:45:35 INFO BlockManagerInfo: Added broadcast_8_piece0 in memory on ip-172-20-5-242.us-east-2.compute.internal:46427 (size: 27.8 KB, free: 424.4 MB)
21/03/21 21:45:35 INFO SparkContext: Created broadcast 8 from broadcast at DAGScheduler.scala:1201
21/03/21 21:45:35 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 7 (MapPartitionsRDD[23] at saveAsTextFile at Driver.scala:28) (first 15 tasks are for partitions Vector(0))
21/03/21 21:45:35 INFO TaskSchedulerImpl: Adding task set 7.0 with 1 tasks
21/03/21 21:45:35 INFO TaskSetManager: Starting task 0.0 in stage 7.0 (TID 43, localhost, executor driver, partition 0, ANY, 7938 bytes)
21/03/21 21:45:35 INFO Executor: Running task 0.0 in stage 7.0 (TID 43)
21/03/21 21:45:35 INFO HadoopMapRedCommitProtocol: Using output committer class org.apache.hadoop.mapred.DirectFileOutputCommitter
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Getting 4 non-empty blocks including 4 local blocks and 0 remote blocks
21/03/21 21:45:35 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms
21/03/21 21:45:35 INFO MultipartUploadOutputStream: close closed:false s3://801198913bucket/FPGrowthOutput1/part-00000
21/03/21 21:45:36 INFO SparkHadoopMapRedUtil: No need to commit output of task because needsTaskCommit=false: attempt_20191027214534_0023_m_000000_0
21/03/21 21:45:36 INFO Executor: Finished task 0.0 in stage 7.0 (TID 43). 1373 bytes result sent to driver
21/03/21 21:45:36 INFO TaskSetManager: Finished task 0.0 in stage 7.0 (TID 43) in 258 ms on localhost (executor driver) (1/1)
21/03/21 21:45:36 INFO TaskSchedulerImpl: Removed TaskSet 7.0, whose tasks have all completed, from pool
21/03/21 21:45:36 INFO DAGScheduler: ResultStage 7 (runJob at SparkHadoopWriter.scala:78) finished in 0.283 s
21/03/21 21:45:36 INFO DAGScheduler: Job 2 finished: runJob at SparkHadoopWriter.scala:78, took 0.937819 s
21/03/21 21:45:36 INFO FileOutputCommitter: File Output Committer Algorithm version is 2
21/03/21 21:45:36 INFO FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: true
21/03/21 21:45:36 INFO DirectFileOutputCommitter: Direct Write: ENABLED
21/03/21 21:45:36 INFO DirectFileOutputCommitter: Nothing to clean up since no temporary files were written.
21/03/21 21:45:36 INFO MultipartUploadOutputStream: close closed:false s3://801198913bucket/FPGrowthOutput1/_SUCCESS
21/03/21 21:45:36 INFO SparkHadoopWriter: Job job_20191027214534_0023 committed.
21/03/21 21:45:36 INFO SparkContext: Invoking stop() from shutdown hook
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 138
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 90
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 188
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 146
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 144

[illegible]

```
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 112
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 162
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 165
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 161
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 111
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 158
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 114
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 194
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 197
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 115
21/03/21 21:45:36 INFO SparkUI: Stopped Spark web UI at http://ip-172-20-
5-242.us-east-2.compute.internal:4040
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 150
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 182
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 98
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 143
21/03/21 21:45:36 INFO ContextCleaner: Cleaned accumulator 154
21/03/21 21:45:36 INFO ContextCleaner: Cleaned shuffle 1
21/03/21 21:45:36 INFO MapOutputTrackerMasterEndpoint:
MapOutputTrackerMasterEndpoint stopped!
21/03/21 21:45:36 INFO MemoryStore: MemoryStore cleared
21/03/21 21:45:36 INFO BlockManager: BlockManager stopped
21/03/21 21:45:36 INFO BlockManagerMaster: BlockManagerMaster stopped
21/03/21 21:45:36 INFO
OutputCommitCoordinator$OutputCommitCoordinatorEndpoint:
OutputCommitCoordinator stopped!
21/03/21 21:45:36 INFO SparkContext: Successfully stopped SparkContext
21/03/21 21:45:36 INFO ShutdownHookManager: Shutdown hook called
21/03/21 21:45:36 INFO ShutdownHookManager: Deleting directory
/mnt/tmp/spark-ab4282d4-3721-48e6-9762-35d60b619cf6
21/03/21 21:45:36 INFO ShutdownHookManager: Deleting directory
/mnt/tmp/spark-3eba5600-85f3-41c0-bf19-3f13ba37e837
[hadoop@ip-172-20-5-242 ~]$
```